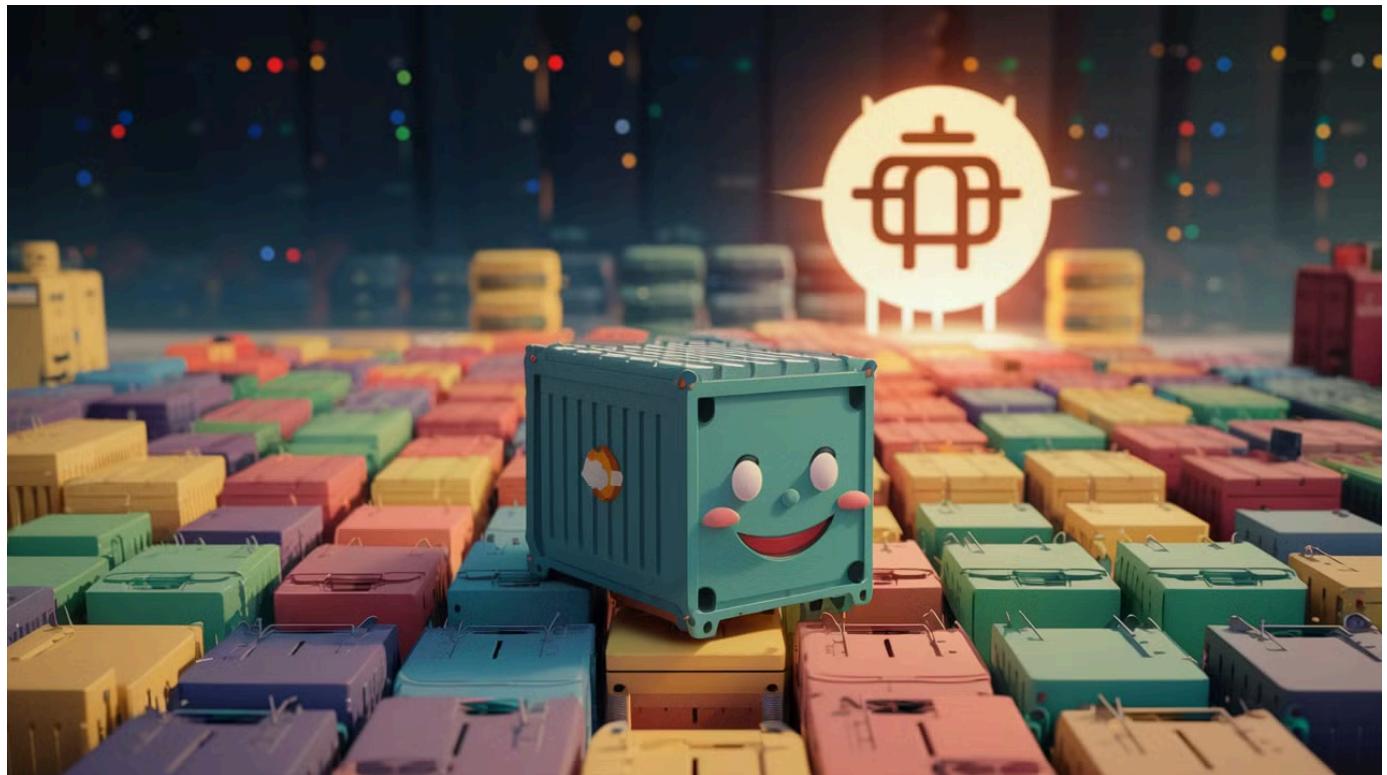




Step 1 - Why Docker?



Docker/containers are important for a few reasons -

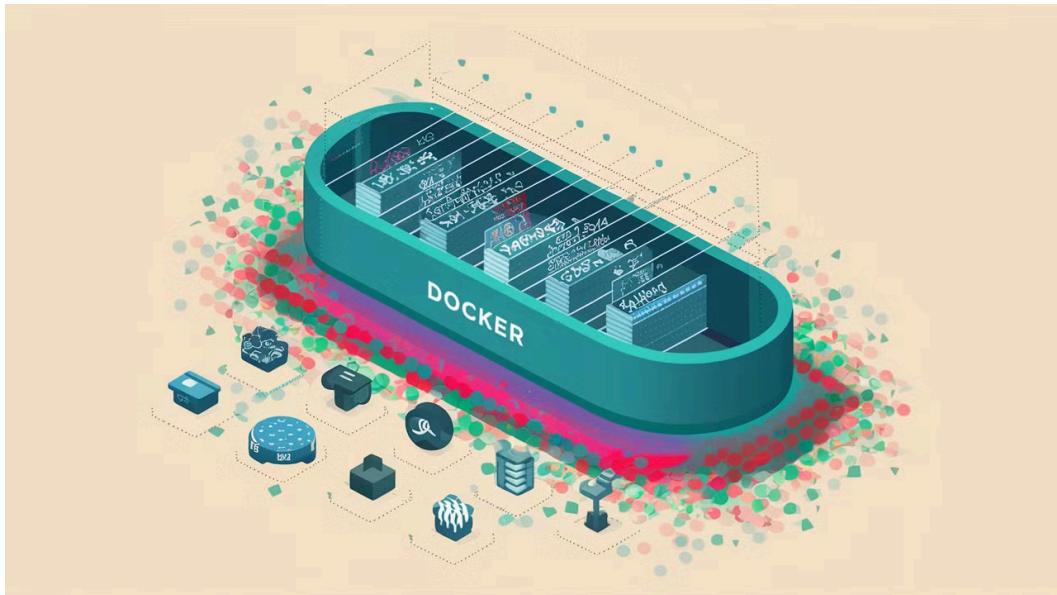
1. Kubernetes/Container orchestration
2. Running processes in isolated environments
3. Starting projects/auxiliary services locally





Step 2 – Containerization

What are containers



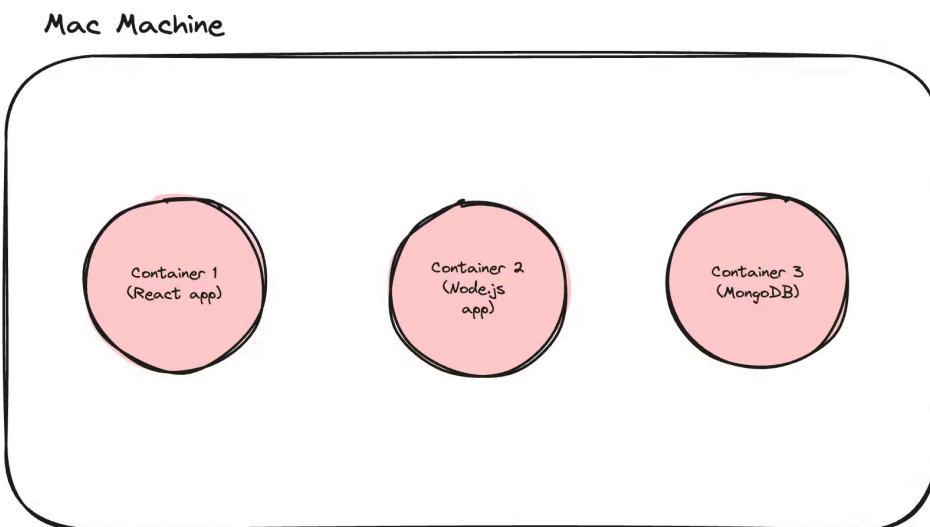
Containers are a way to package and distribute software applications in a way that makes them easy to deploy and run consistently across different environments. They allow you to package an application, along with all its dependencies and libraries, into a single unit that can be run on any machine with a container runtime, such as Docker.

Why containers



1. Everyone has different Operating systems
2. Steps to run a project can vary based on OS
3. Extremely harder to keep track of dependencies as project grows

Benefits of using containers



1. Let you describe your **configuration** in a single file
2. Can run in isolated environments
3. Makes Local setup of OS projects a breeze

It's easy

References

Docker End to End 1 of 22

Step 3 – History of Docker



Docker is a YC backed company, started in ~2014

They envisioned a world where containers would become mainstream and people would deploy their applications using them

That is mostly true today

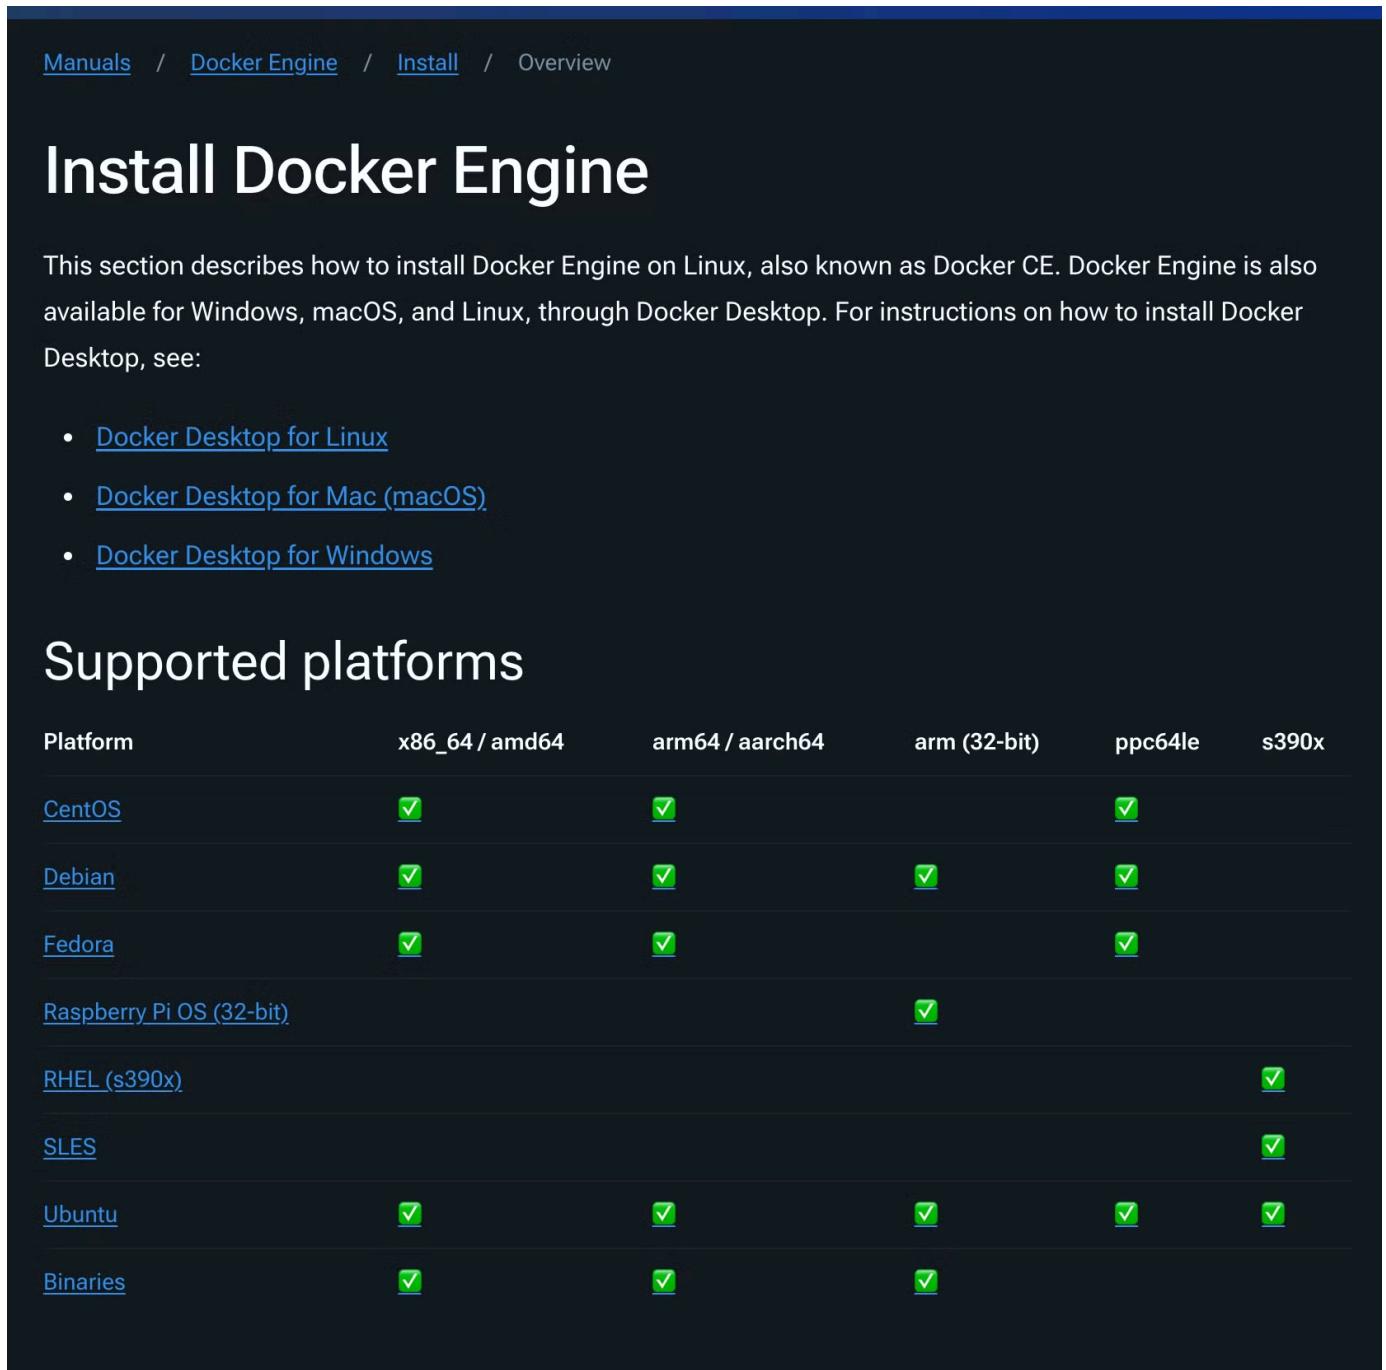
Most projects that you open on Github will/should have docker files in them
(a way to create docker containers)

Ref - <https://www.ycombinator.com/blog/solomon-hykes-docker-dotcloud-interview/>

 Docker End to End 1 of 22

Step 4 – Installing docker

<https://docs.docker.com/engine/install/>



The screenshot shows the Docker Engine installation documentation. The top navigation bar includes links for Manuals, Docker Engine, Install, and Overview. The main title is "Install Docker Engine". A descriptive paragraph explains that this section covers Linux installation, mentioning Docker CE, and provides links for Docker Desktop on various platforms. Below this is a section titled "Supported platforms" with a table.

Platform	x86_64 / amd64	arm64 / aarch64	arm (32-bit)	ppc64le	s390x
CentOS	✓	✓		✓	
Debian	✓	✓	✓	✓	
Fedora	✓	✓		✓	
Raspberry Pi OS (32-bit)			✓		
RHEL (s390x)				✓	
SLES				✓	
Ubuntu	✓	✓	✓	✓	✓
Binaries	✓	✓	✓		

N

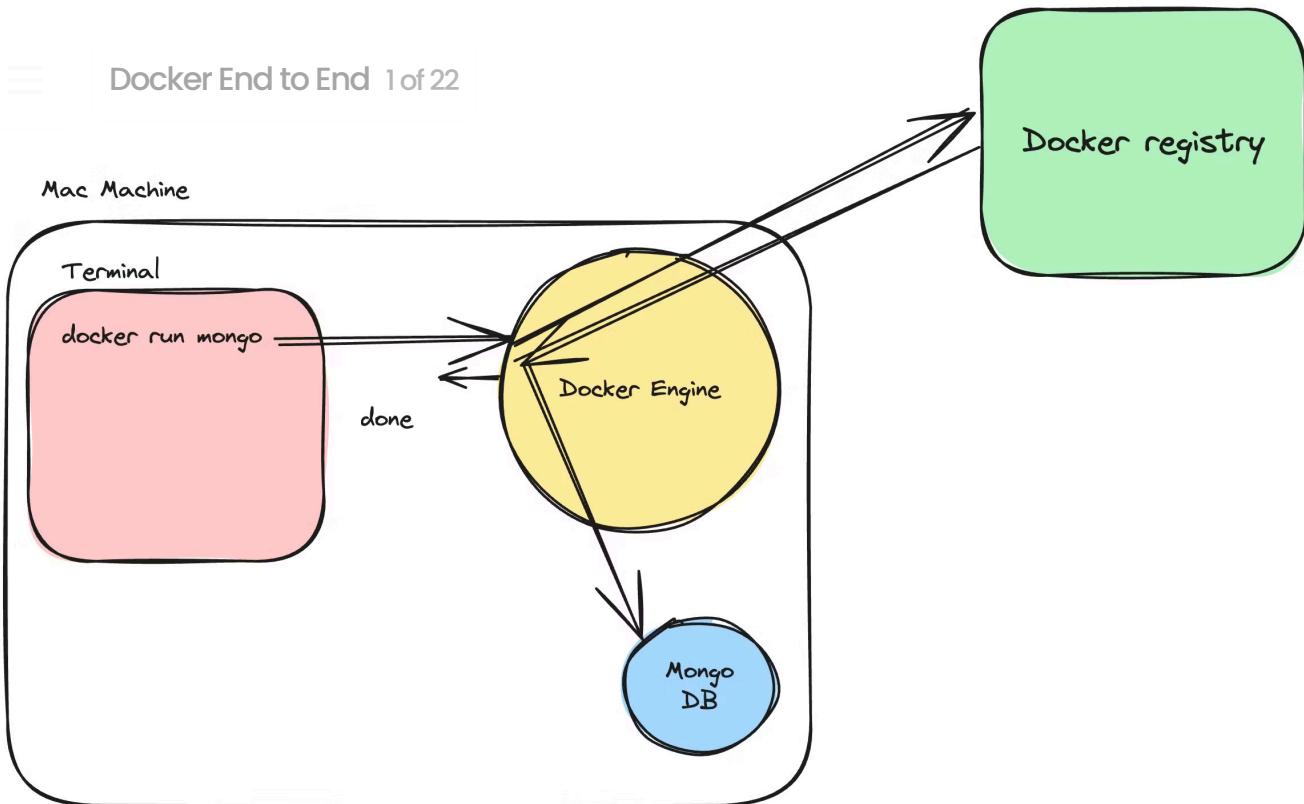
cli locally –

```
→ ~ docker
  Docker End to End 1 of 22 COMMAND
A self-sufficient runtime for containers

Common Commands:
run          Create and run a new container from an image
exec         Execute a command in a running container
ps           List containers
build        Build an image from a Dockerfile
pull         Download an image from a registry
push         Upload an image to a registry
images       List images
login        Log in to a registry
logout       Log out from a registry
search       Search Docker Hub for images
version      Show the Docker version information
info         Display system-wide information
```

Step 5 – Inside docker

Docker End to End 1 of 22



As an application/full stack developer, you need to be comfortable with the following terminologies –

1. Docker Engine
2. Docker CLI – Command line interface
3. Docker registry

1. Docker Engine

Docker Engine is an open-source **containerization** technology that allows developers to package applications into **container**

Containers are standardized executable components combining application source code with the operating system (os) libraries and dependencies required to run that code in any environment.

2. Docker CLI

The command line interface lets you talk to the **docker engine** and lets you ~~start/stop/inspect containers~~

```
'Docker run -d -p 27017:27017 mongo'
```

Docker End to End 1 of 22

 Docker cli is not the only way to talk to a docker engine. You can hit the docker REST API to do the same things

Step 6 – Images vs containers

Docker Image

A Docker image is a lightweight, standalone, executable package that includes everything needed to run a piece of software, including the code, a runtime, libraries, environment variables, and config files.

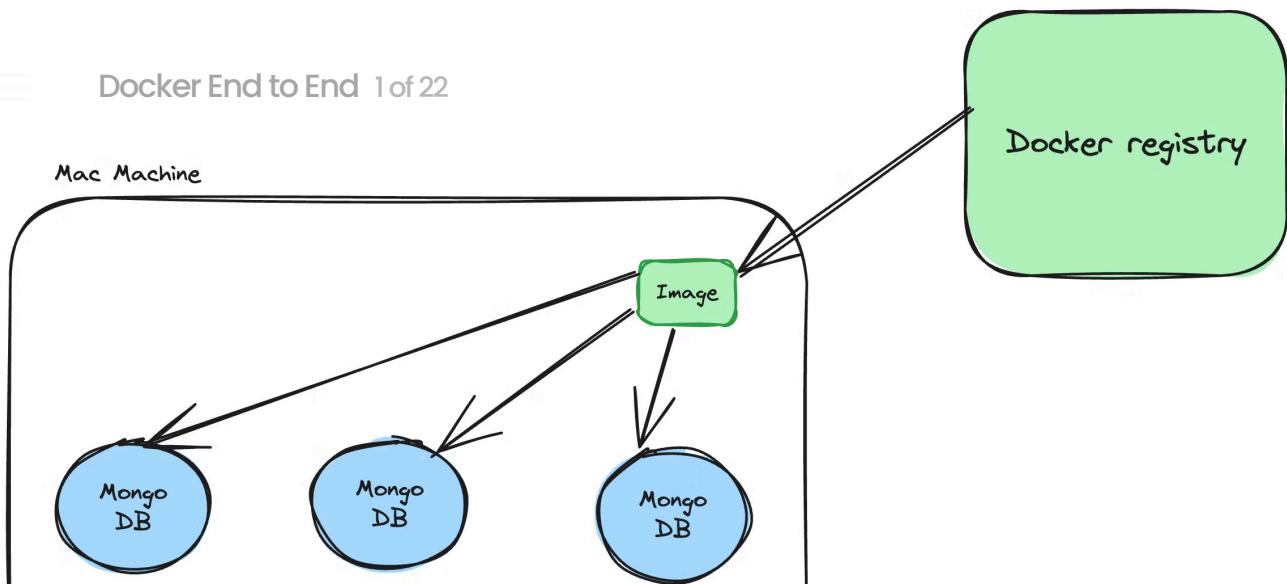
 A good mental model for an image is Your codebase on github

Docker Container

A container is a running instance of an image. It encapsulates the application or service and its dependencies, running in an isolated environment.

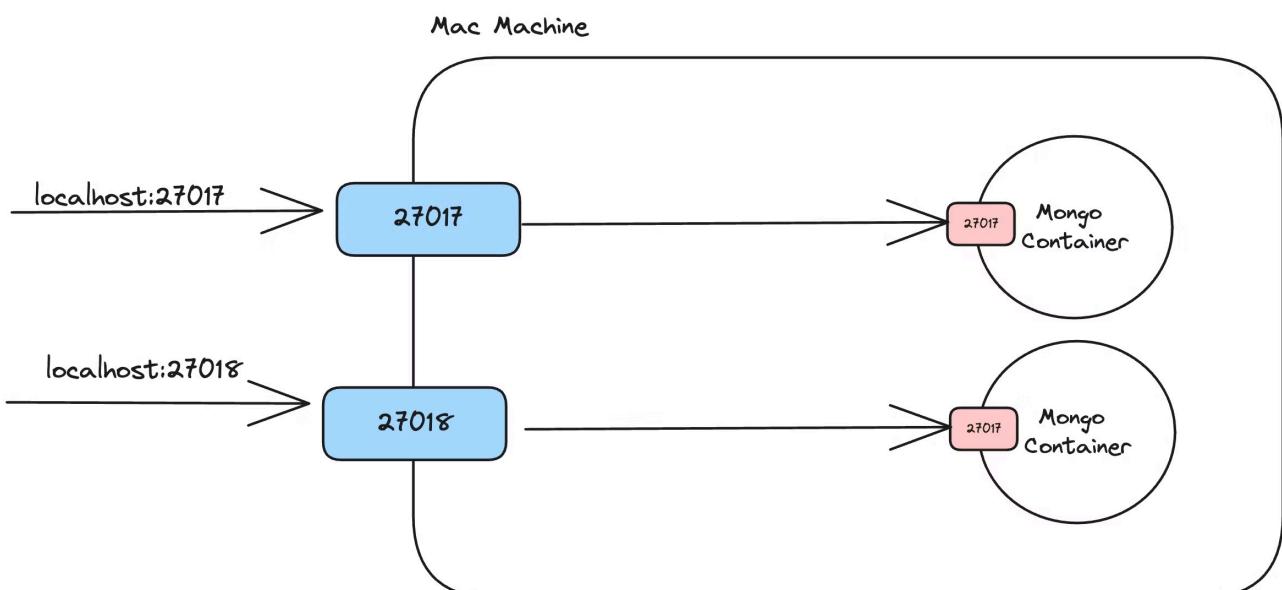
 A good mental model for a container is when you run node index.js on your machine from some source code you got from github

Docker End to End 1 of 22



Step 7 – Port mapping

```
docker run -d -p 27018:27017 mongo
```





Step 8 – Common docker commands

1. docker images
2. docker ps
3. docker run
4. docker build

1. docker images

Shows you all the images that you have on your machine

2. docker ps

Shows you all the containers you are running on your machine

3. docker run

L

1 -n ⇒ let's you create a port mapping
≡ Docker End to End 1 of 22 :atched mode

4. docker build

Lets you build an image. We will see this after we understand how to create your own **Dockerfile**

5. docker push

Lets you push your image to a registry

6. Extra commands

1. docker kill
2. docker exec

Step 9 – Dockerfile

What is a Dockerfile

If you want to create an image from your own code, that you can push to **dockerhub**, you need to create a **Dockerfile** for your application.

A Dockerfile is a text document that contains all the commands a user could call on the command line to create an image.

How to write a Dockerfile

A dockerfile has 2 parts

☰ ; Docker End to End 1 of 22

2. Bunch of commands that you run on the base image (to install dependencies like Node.js)

Let's write our own Dockerfile

Let's try to containerise this backend app - <https://github.com/100xdevs-cohort-2/week-15-live-1>

```
📄 Dockerfile
1  FROM node:16-alpine → Base Image
2
3  WORKDIR /app → Working directory
4
5  COPY . . → Copy over files
6
7  RUN npm install → Run Commands to build the code
8  RUN npm run build
9
10 EXPOSE 3000 → Expose ports
11
12 CMD ["node", "dist/index.js"] → Final command that runs when running the
13
  container
```

▼ Solution

```
FROM node:20
```



```
WORKDIR /app
```

```
COPY ..
```

```
RUN npm install
```

```
RUN npx prisma generate
```

```
RUN npm run build
```

```
EXPOSE 3000
```

```
CMD ["node", "dist/index.js"]
```

Common commands

- **WORKDIR** : Sets the working directory for any **RUN** , **CMD** , **ENTRYPOINT** ,
☰ ↗ Docker End to End 1 of 22 low it.
- **RUN** : Executes any commands in a new layer on top of the current image and commits the results.
- **CMD** : Provides defaults for executing a container. There can only be one CMD instruction in a Dockerfile.
- **EXPOSE** : Informs Docker that the container listens on the specified network ports at runtime.
- **ENV** : Sets the environment variable.
- **COPY** : Allow files from the Docker host to be added to the Docker image

<https://github.com/100xdevs-cohort-2/week-15-live-1>

Step 10 – Building images

Now that you have a dockerfile in your project, try building a **docker image** from it

docker build -t image_name .



Now if you try to look at your images, you should notice a new image created

docker images





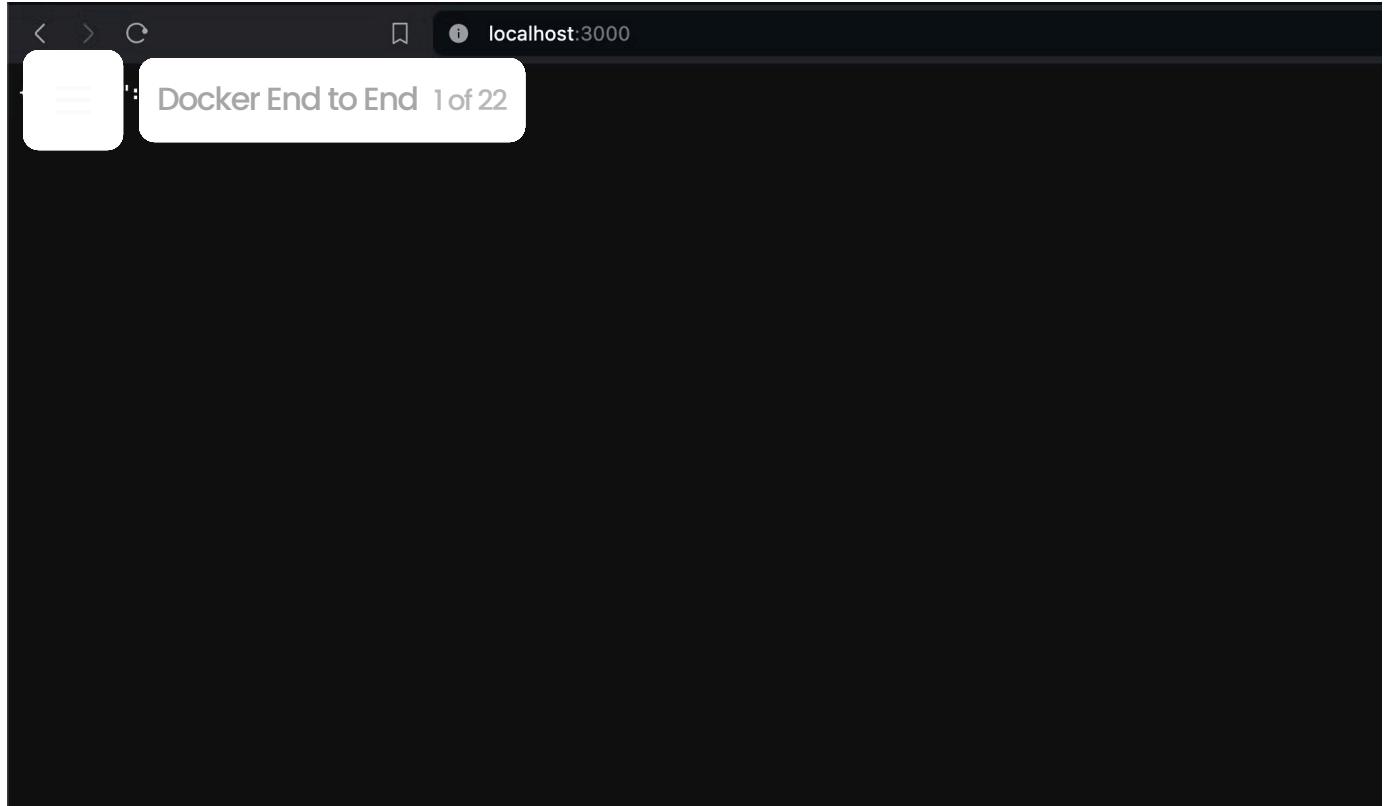
\r'--- that node_modules don't get copied over
Docker End to End 1 of 22

Step 11 – Running images

`docker run -p 3000:3000 image_name`



Try visiting `localhost:3000`



Step 12 – Passing in env variables

```
docker run -p 3000:3000 -e DATABASE_URL="postgres://avnadmin:AVNS_EeDil
```



The `-e` argument lets you send in environment variables to your node.js app



Step 13 – More commands

1. docker kill – to kill a container
2. docker exec – to execute a command inside a container

Examples

1. List all contents of a container folder

```
docker exec <container_name_or_id> ls /path/to/directory
```



1. Running an Interactive Shell

```
> /bin/bash
```





Step 14 – Pushing to dockerhub

Once you've created your image, you can push it to [dockerhub](#) to share it with the world.

1. Signup to [dockerhub](#)
2. Create a new repository

Docker End to End 1 of 22

Search by repository name All Content

Create repository

100xdevs / runner Contains: Image Last pushed: 22 days ago	Security unknown	0	55	Public
100xdevs / test-during-video Contains: Image Last pushed: 5 months ago	Security unknown	0	29	Public
100xdevs / testing Contains: Image Last pushed: 5 months ago	Security unknown	0	8	Public
100xdevs / simple-node-app Contains: No content Created: 5 months ago	Security unknown	0	0	Public

Create An Organization

Create and manage users and grant access to your

3. Login to docker cli

1. docker login

2. you might have to create an access token -

<https://docs.docker.com/security/for-developers/access-tokens/>

4. Push to the repository

docker push your_username/your_reponame:tagname

Step 15 – Layers in Docker

In Docker, layers are a fundamental part of the image architecture that allows Docker to be efficient, fast, and portable. A Docker image is essentially built up from a series of layers, each representing a set of differences from the previous layer.



Docker End to End 1 of 22

How layers are made -

- 1. Base Layer:** The starting point of an image, typically an operating system (OS) like Ubuntu, Alpine, or any other base image specified in a Dockerfile.
- 2. Instruction Layers:** Each command in a Dockerfile creates a new layer in the image. These include instructions like `RUN`, `COPY`, which modify the filesystem by installing packages, copying files from the host to the container, or making other changes. Each of these modifications creates a new layer on top of the base layer.
- 3. Reusable & Shareable:** Layers are cached and reusable across different images, which makes building and sharing images more efficient. If multiple images are built from the same base image or share common instructions, they can reuse the same layers, reducing storage space and speeding up image downloads and builds.
- 4. Immutable:** Once a layer is created, it cannot be changed. If a change is made, Docker creates a new layer that captures the difference. This immutability is key to Docker's reliability and performance, as unchanged layers can be shared across images and containers.

Step 16 – Layers practically

For a simple Node.js app - <https://github.com/100xdevs-cohort-2/week-15-live-2>

Dockerfile

```
FROM node:20
WORKDIR /usr/src/app
COPY . .
RUN npm install
RUN npm run build
RUN npx prisma generate
EXPOSE 3000
CMD ["node", "dist/index.js", ]
```

Logs

```
● → week-15-live-1 git:(main) ✘ docker build -t hkirat-backend-cool-app .
[+] Building 12.8s (12/12) FINISHED                                            docker:desktop-linux
  => [internal] load .dockerignore                                         0.0s
  => => transferring context: 57B                                           0.0s
  => [internal] load build definition from Dockerfile                   0.0s
  => => transferring dockerfile: 189B                                      0.0s
  => [internal] load metadata for docker.io/library/node:20            3.7s
  => [auth] library/node:pull token for registry-1.docker.io           0.0s
  => CACHED [1/6] FROM docker.io/library/node:20@sha256:f3299f16246c71ab8b304d6745bb4059f 0.0s
  => [internal] load build context                                         0.0s
  => => transferring context: 7.55kB                                     0.0s
  => [2/6] WORKDIR /usr/src/app                                         0.0s
  => [3/6] COPY . .                                                 0.0s
  => [4/6] RUN npm install                                              6.0s
  => [5/6] RUN npm run build                                             1.3s
  => [6/6] RUN npx prisma generate                                       1.1s
  => exporting to image                                                 0.5s
                                                               0.4s
                                                               e075ad09dedebbfdafe8353364cf35bf8dafd8 0.0s
                                                               -cool-app                                         0.0s
```

 Docker End to End 1 of 22

Observations -

1. Base image creates the first layer
2. Each `RUN`, `COPY`, `WORKDIR` command creates a new layer
3. Layers can get re-used across docker builds (notice `CACHED` in 1/6)

Step 17 – Why layers?

If you change your Dockerfile, layers can get re-used based on where the change was made



If a layer changes, all subsequent layers also change

Case 1 – You change your source code

```
📦 Dockerfile
1  FROM node:20
2
3  WORKDIR /usr/src/app
4
5  COPY . .
6
7  RUN npm install
8  RUN npm run build
9  RUN npx prisma generate
10
11
12 EXPOSE 3000
13
14 CMD ["node", "dist/index.js", ]
```

Layer 1 - Cached
Layer 2 - Cached
Layer 3 - Changed
Layer 4 - Changed
Layer 5 - Changed
Layer 6 - Changed

```

● → week-15-live-1 git:(main) ✘ docker build -t hkirat-backend-cool-app .
[+] Building 8.7s (11/11) FINISHED
  => [internal] load build definition from Dockerfile
  => [internal] load .dockerignore
  => => transferring dockerfile: 189B
  => [internal] load .dockerignore
  => => transferring context: 57B
  => [internal] load metadata for docker.io/library/node:20
  => [1/6] FROM docker.io/library/node:20@sha256:f3299f16246c71ab8b304d6745bb4059fa9283e8d025972e28436a9f9b36ed24
  => [internal] load build context
  => => transferring context: 3.51kB
  => CACHED [2/6] WORKDIR /usr/src/app
  => [3/6] COPY .
  => [4/6] RUN npm install
  => [5/6] RUN npm run build
  => [6/6] RUN npx prisma generate
  => exporting to image
  => => exporting layers
  => => writing image sha256:e0a88f16c2d11e317070c51ec298e61325c0ef4a2214d2c11a2cc063d4b5e338
  => => naming to docker.io/library/hkirat-backend-cool-app

View build details: docker-desktop://dashboard/build/desktop-linux/desktop-linux/yimvu9plvmxbypwatyvd4mdg

What's Next?
  View a summary of image vulnerabilities and recommendations → docker scout quickview

```

Is still cached, even though it doesn't say it

Case 2 – You change the package.json file (added a dependency)

```

Dockerfile
1  FROM node:20
2
3  WORKDIR /usr/src/app
4
5  COPY .
6
7  RUN npm install
8  RUN npm run build
9  RUN npx prisma generate
10
11
12 EXPOSE 3000
13
14 CMD ["node", "dist/index.js", ]

```

Layer 1 - Cached

Layer 2 - Cached

Layer 3 - Changed

Layer 4 - Changed

Layer 5 - Changed

Layer 6 - Changed

▼ Logs

```

● → week-15-live-1 git:(main) ✘ docker build -t hkirat-backend-cool-app .
[+] Building 8.7s (11/11) FINISHED
  => [internal] load build definition from Dockerfile
  => => transferring dockerfile: 189B
  => [internal] load .dockerignore
  => => transferring context: 57B
  => [internal] load metadata for docker.io/library/node:20
  => [1/6] FROM docker.io/library/node:20@sha256:f3299f16246c71ab8b304d6745bb4059fa9283e8d025972e28436a9f9b36ed24
  => [internal] load build context
  => => transferring context: 3.51kB
  => CACHED [2/6] WORKDIR /usr/src/app
  => [3/6] COPY .
  => [4/6] RUN npm install
  => [5/6] RUN npm run build
  => [6/6] RUN npx prisma generate
  => exporting to image
  => => exporting layers
  => => writing image sha256:e0a88f16c2d11e317070c51ec298e61325c0ef4a2214d2c11a2cc063d4b5e338
  => => naming to docker.io/library/hkirat-backend-cool-app

View build details: docker-desktop://dashboard/build/desktop-linux/desktop-linux/yimvu9plvmxbypwatyvd4mdg

What's Next?
  View a summary of image vulnerabilities and recommendations → docker scout quickview

```

Is still cached, even though it doesn't say it

How often in a project do you think `dependencies change` ?

F ⌂ Docker End to End 1 of 22 3/11 layer need to change?

Wouldn't it be nice if we could `cache` the `npm install` step considering dependencies don't change often?

Step 18 – Optimising Dockerfile

What if we change the Dockerfile a bit -

```
⌚ Dockerfile
1  FROM node:20
2
3  WORKDIR /usr/src/app
4
5  COPY package* .
6  COPY ./prisma .
7
8  RUN npm install
9  RUN npx prisma generate
10
11 COPY . .
12
13 EXPOSE 3000
14
15 CMD ["node", "dist/index.js", ]
```

▼ Dockerfile

FROM node:20



```
COPY package* .

  Docker End to End 1 of 22
  RUN npm install
  RUN npx prisma generate

```

```
COPY ..
```

```
EXPOSE 3000
```

```
CMD ["node", "dist/index.js", ]
```

1. We first copy over only the things that `npm install` and `npx prisma generate` need
2. Then we run these scripts
3. Then we copy over the rest of the source code

Case 1 – You change your source code (but nothing in package.json/prisma)

```


The Dockerfile is visualized as a stack of layers. Each line of the Dockerfile is preceded by a line number and followed by its corresponding layer status.



```

1 FROM node:20
2
3 WORKDIR /usr/src/app
4
5 COPY package* .
6 COPY ./prisma .
7
8 RUN npm install
9 RUN npx prisma generate
10
11 COPY . .
12
13 EXPOSE 3000
14
15 CMD ["node", "dist/index.js",]

```



- Layer 1 - Cached: FROM node:20
- Layer 2 - Cached: WORKDIR /usr/src/app
- Layer 3 - Cached: COPY package* .
- Layer 4 - Cached: COPY ./prisma .
- Layer 5 - Cached: RUN npm install
- Layer 6 - Cached: RUN npx prisma generate
- Layer 7 - Uncached: COPY . . (A green arrow points to this line from the left margin.)
- Layer 8 - Uncached: EXPOSE 3000
- Layer 9 - Uncached: CMD ["node", "dist/index.js", ]

```

Case 2 – You change the package.json file (added a command)

The screenshot shows a Dockerfile editor interface. On the left, there's a sidebar with a file icon and the text "Dockerfile". The main area contains a Dockerfile with numbered lines from 1 to 15. Lines 1 through 14 are grouped by a blue vertical bar on the left, while line 15 is grouped by a green vertical bar. To the right of the lines, there are two columns of text: "Layer 1 - Cached" (green) and "Layer 2 - Cached" (blue). Lines 1-4 are in the first column, and lines 5-14 are in the second. Lines 15-17 are grouped by a red vertical bar on the left, with "Layer 3 - Uncached" (red) above them. Lines 18-20 are grouped by a blue vertical bar on the left, with "Layer 4 - Uncached" (blue) above them. Lines 21-23 are grouped by a red vertical bar on the left, with "Layer 5 - Uncached" (red) above them. Lines 24-26 are grouped by a blue vertical bar on the left, with "Layer 6 - Uncached" (blue) above them. Lines 27-29 are grouped by a red vertical bar on the left, with "Layer 7 - Uncached" (red) above them. Lines 30-32 are grouped by a blue vertical bar on the left, with "Layer 8 - Uncached" (blue) above them.

```
1 WORKDIR /usr/src/app  
2  
3 COPY package*.json .  
4 COPY ./prisma .  
5  
6 RUN npm install  
7 RUN npx prisma generate  
8  
9 COPY . .  
10  
11 EXPOSE 3000  
12  
13 CMD ["node", "dist/index.js", ]
```

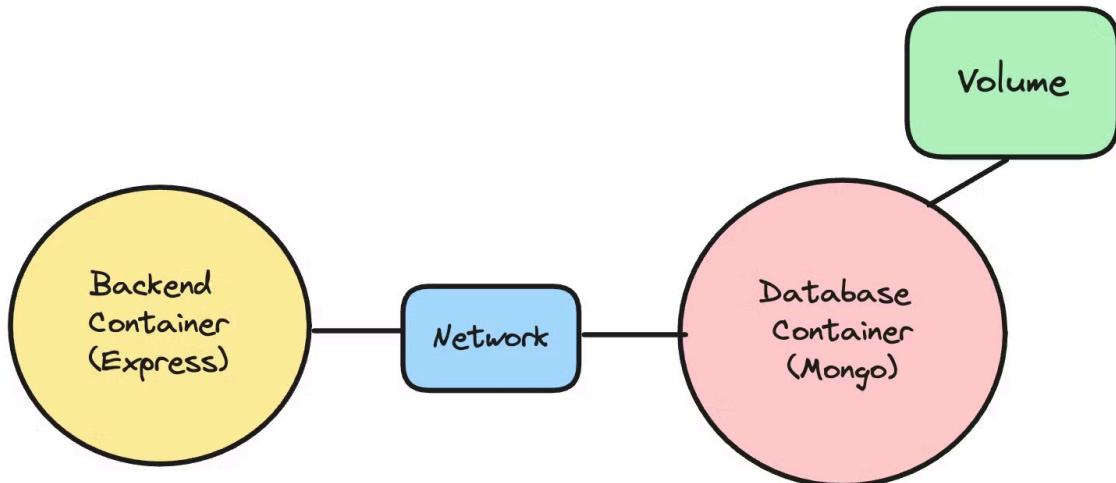
Step 19 – Networks and volumes

Networks and volumes are concepts that become important when you have multiple containers running in which you

1. Need to persist data across docker restarts
2. Need to allow containers to talk to each other

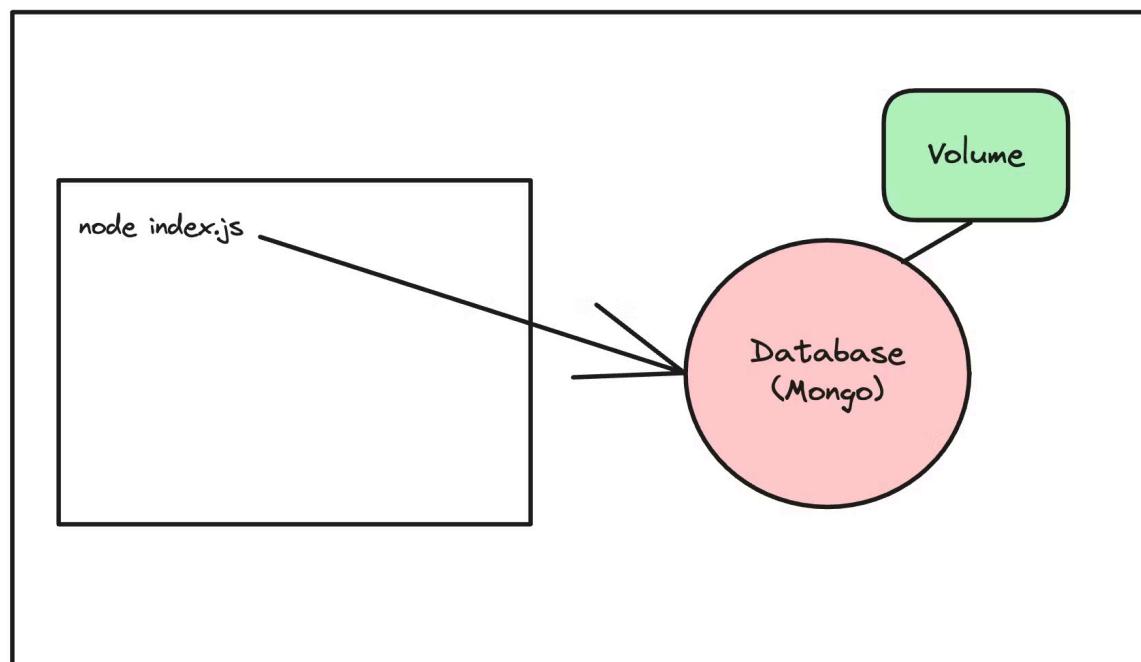
Mac

Docker End to End 1 of 22



We didn't need `networks` until now because when we started the `mongo` container, it was being accessed by a Node.js process running directly on the machine

Mac



Docker End to End 1 of 22 plumes

If you restart a `mongo` docker container, you will notice that your data goes away.

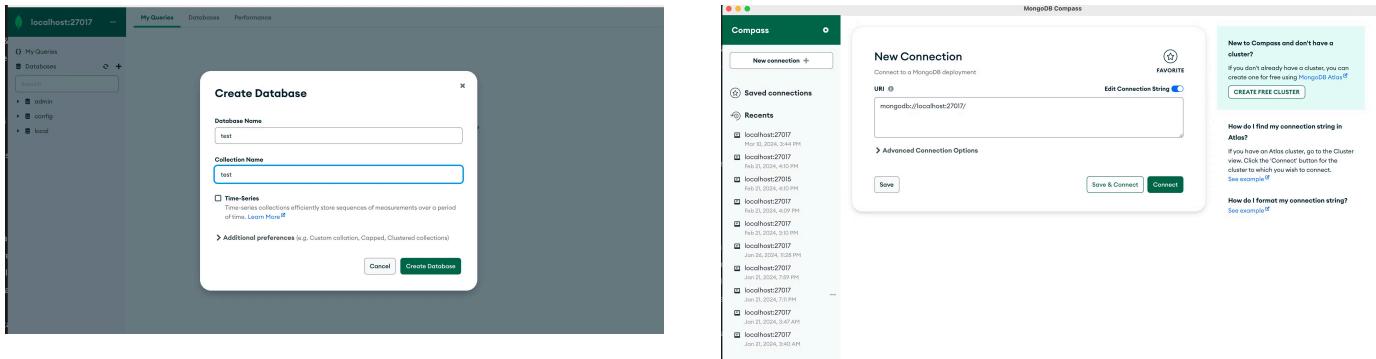
This is because docker containers are **transitory** (they don't retain data across restarts)

Without volumes

1. Start a mongo container locally

```
docker run -p 27017:27017 -d mongo
```

1. Open it in MongoDB Compass and add some data to it



1. Kill the container

```
docker kill <container_id>
```

1. Restart the container

```
docker run -p 27017:27017 -d mongo
```

1. Try to explore the database in Compass and check if the data has persisted (it wouldn't)

1 Create a volume

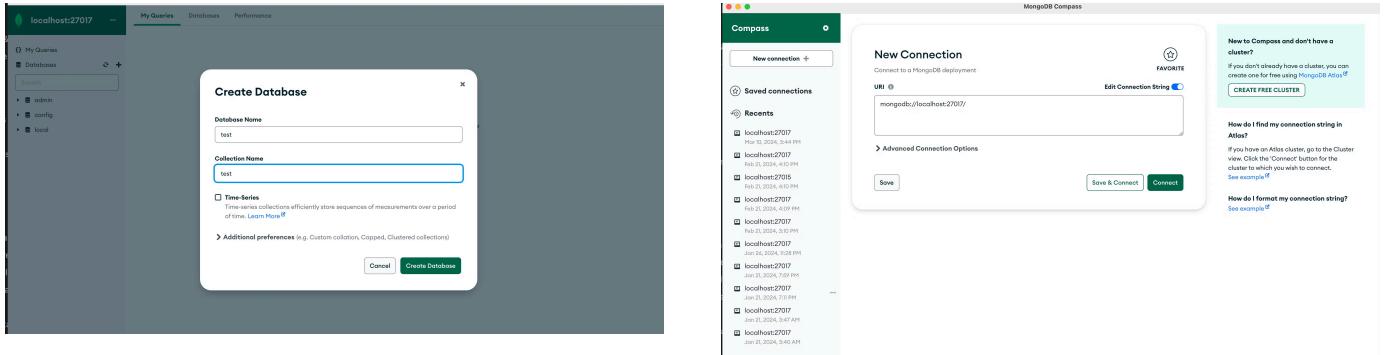
Docker End to End 1 of 22

[CREATE VOLUME](#) [CREATE VOLUME_DATABASE](#)

1. Mount the folder in `mongo` which actually stores the data to this volume

```
docker run -v volume_database:/data/db -p 27017:27017 mongo
```

1. Open it in MongoDB Compass and add some data to it



1. Kill the container

```
docker kill <container_id>
```

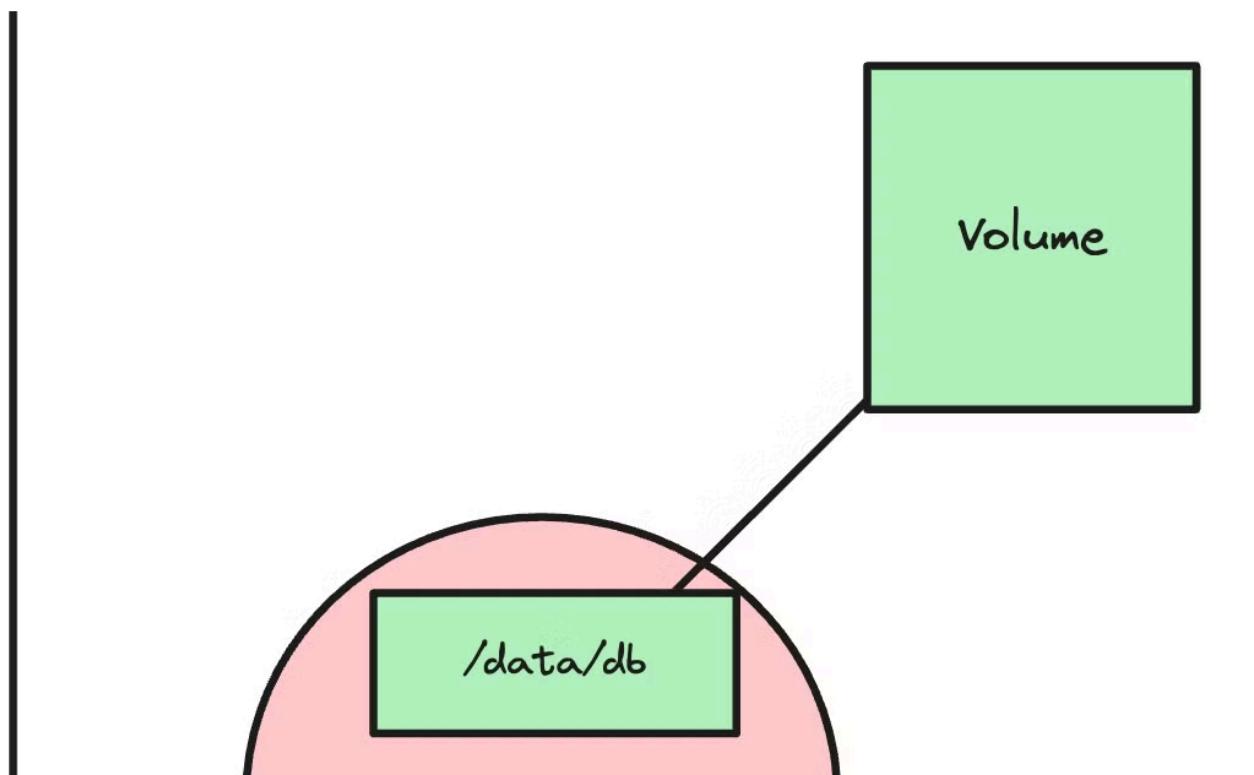
1. Restart the container

```
docker run -v volume_database:/data/db -p 27017:27017 mongo
```

1. Try to explore the database in Compass and check if the data has persisted (it will!)

Mac

Docker End to End 1 of 22



Step 21 – Network

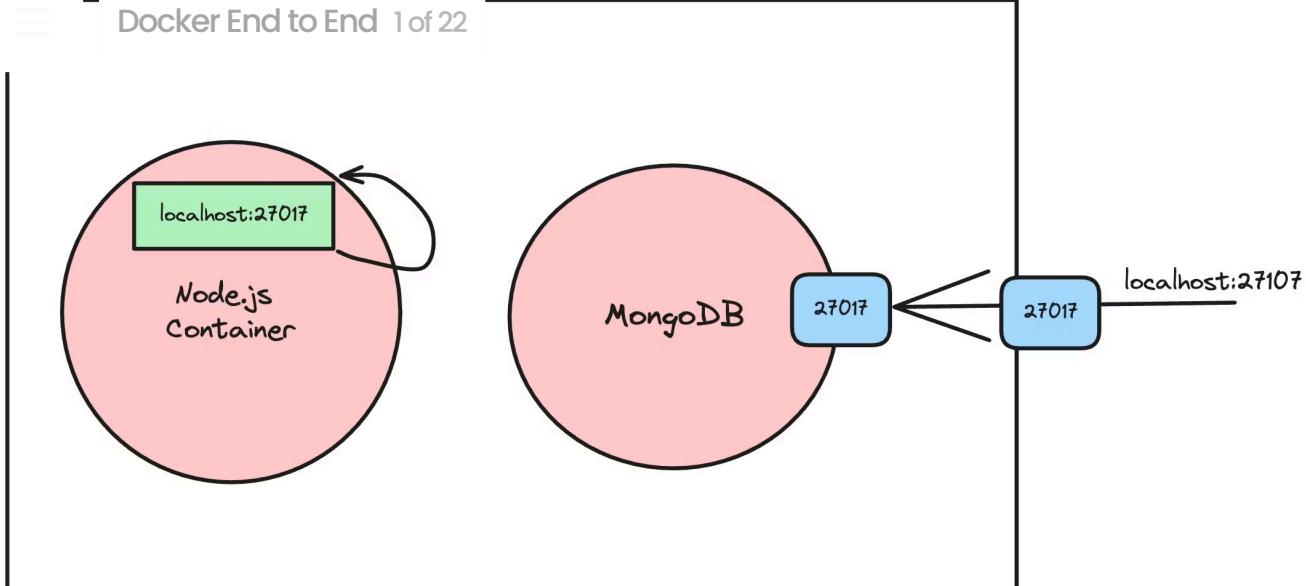
In Docker, a network is a powerful feature that allows containers to communicate with each other and with the outside world.

Docker containers can't talk to each other by default.

`localhost` on a docker container means **it's own network** and not the network of the **host machine**

Map Machine

Docker End to End 1 of 22



How to make containers talk to each other?

Attach them to the same network

1. Clone the repo - <https://github.com/100xdevs-cohort-2/week-15-live-2.2>
2. Build the image

```
docker build -t image_tag .
```



1. Create a network

```
docker network create my_custom_network
```



1. Start the **backend process** with the **network** attached to it

```
docker run -d -p 3000:3000 --name backend --network my_custom_network
```



1. Start mongo on the same network

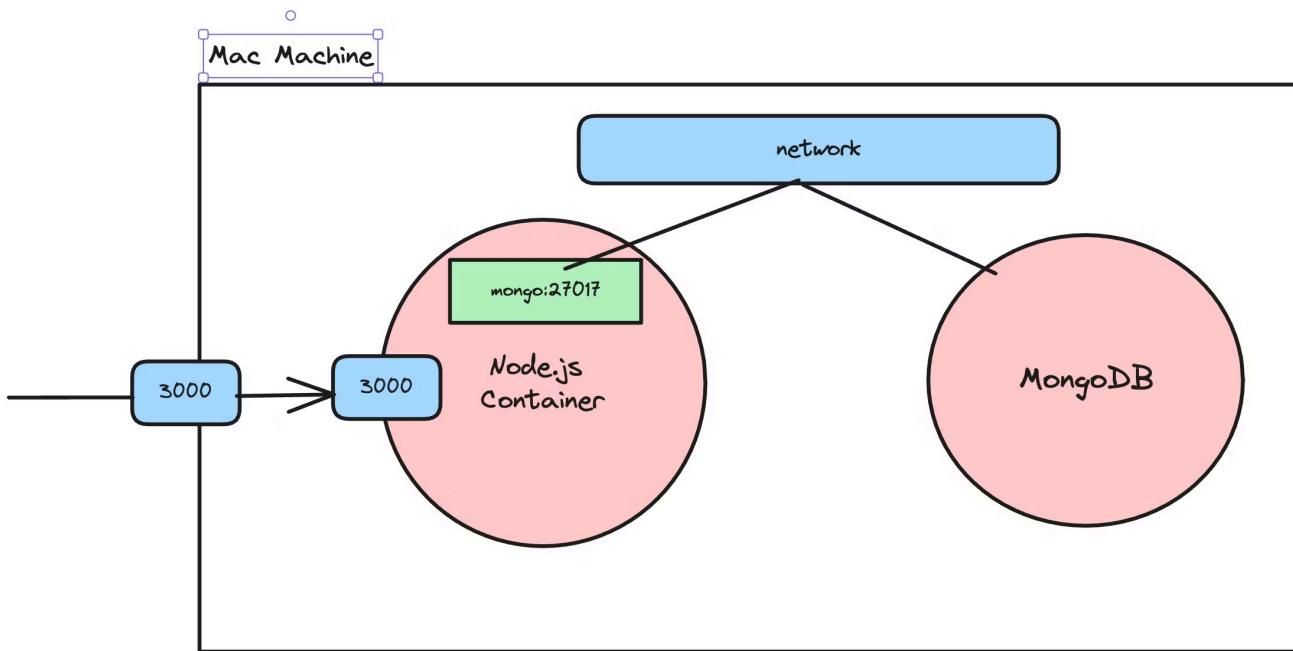
```
docker run -d -v volume_database:/data/db --name mongo --network my_
```



1. Check the logs to ensure the db connection is successful



- Trv to visit an endpoint and ensure you are able to talk to the database
 - Docker End to End 1 of 22
 - Give the port mapping for mongo since you don't necessarily need it exposed on your machine



Types of networks

- Bridge:** The default network driver for containers. When you run a container without specifying a network, it's attached to a bridge network. It provides a private internal network on the host machine, and containers on the same bridge network can communicate with each other.
- Host:** Removes network isolation between the container and the Docker host, and uses the host's networking directly. This is useful for services

Step 22 - docker-compose

Docker Compose is a tool designed to help you define and run multi-container Docker applications. Instead of running multiple Docker commands, you use a YAML file to configure your application's services, networks, and volumes. Then, with a single command, you can start, stop, and manage all the services in your application.

single command, you can create and start all the services from your

C ⌂ Docker End to End 1 of 22

```
docker-compose.yaml
1 version: '3.8'                                Version of docker-compose
2 services:
3   mongodb:
4     image: mongo
5     container_name: mongodb
6     ports:
7       - "27017:27017"
8     volumes:
9       - mongodb_data:/data/db
10
11  backend22:
12    image: backend22
13    container_name: backend_app
14    depends_on:
15      - mongodb
16    ports:
17      - "3000:3000"
18    environment:
19      MONGO_URL: "mongodb://mongodb:27017"
20
21    volumes:
22      - mongodb_data: Volumes to create
```

Before docker-compose

- Create a network

```
docker network create my_custom_network
```



- Create a volume

```
docker volume create volume_database
```



- Start mongo container

```
docker run -d -v volume_database:/data/db --name mongo --network my_
```



```
'docker run -d -p 3000:3000 --name backend --network my_custom_network'
```

Docker End to End 1 of 22

After docker-compose

1. Install docker-compose - <https://docs.docker.com/compose/install/>
2. Create a **yaml** file describing all your containers and volumes (by default all containers in a docker-compose run on the same network)

▼ Solution

```
version: '3.8'  
services:  
  mongodb:  
    image: mongo  
    container_name: mongodb  
    ports:  
      - "27017:27017"  
    volumes:  
      - mongodb_data:/data/db  
  
  backend22:  
    image: backend  
    container_name: backend_app  
    depends_on:  
      - mongodb  
    ports:  
      - "3000:3000"  
    environment:  
      MONGO_URL: "mongodb://mongodb:27017"  
  
    volumes:  
      mongodb_data:
```

1. Start the compose

docker-compose up

1. Stop everything (including volumes)



Docker End to End 1 of 22