



# What we're discussing

## Server

1. Cluster module and horizontal scaling ✓
2. Capacity Estimation, ASGs and Vertical scaling ✓
3. Load balancers ✓

## Database

1. Indexing ✓
2. Normalization
3. Sharding

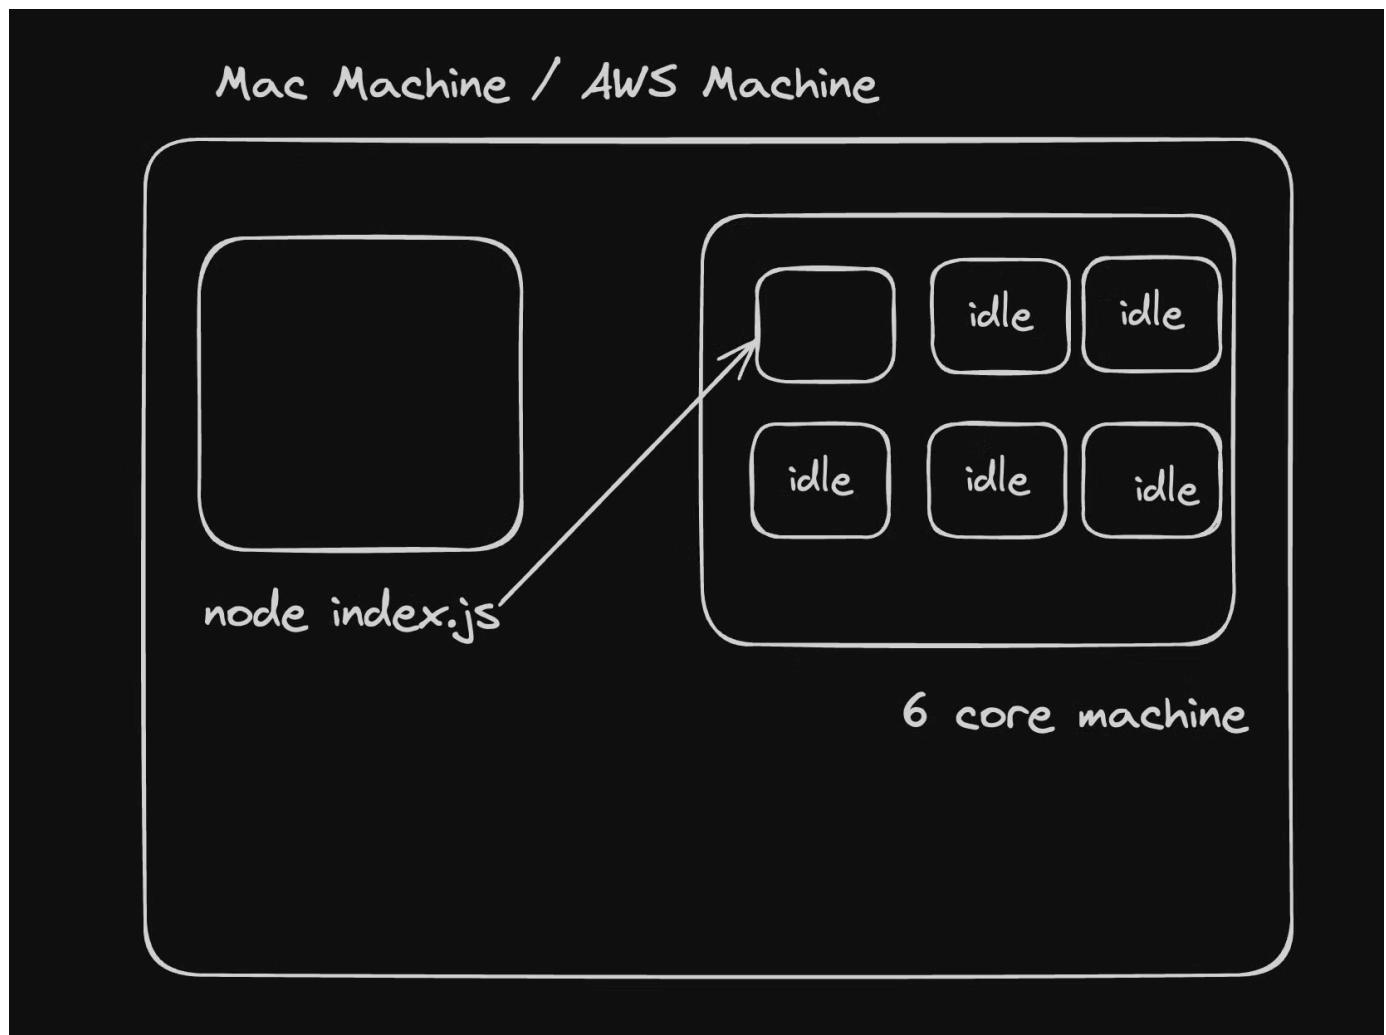




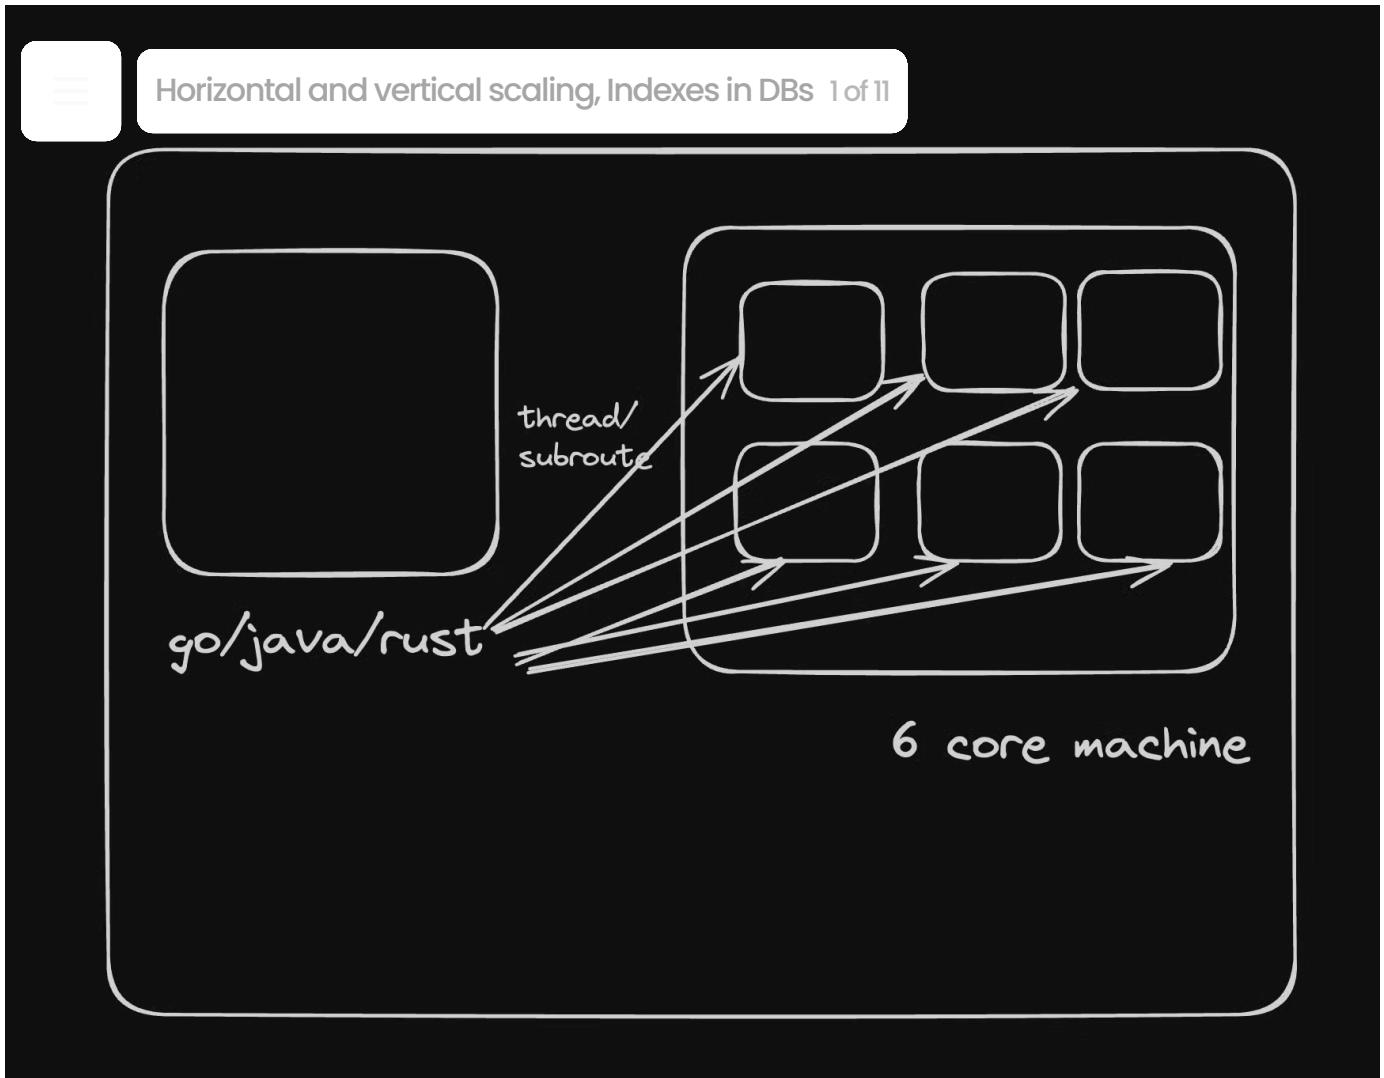
# Vertical scaling

Vertical scaling means increasing the size of your machine to support more load

## Single threaded languages



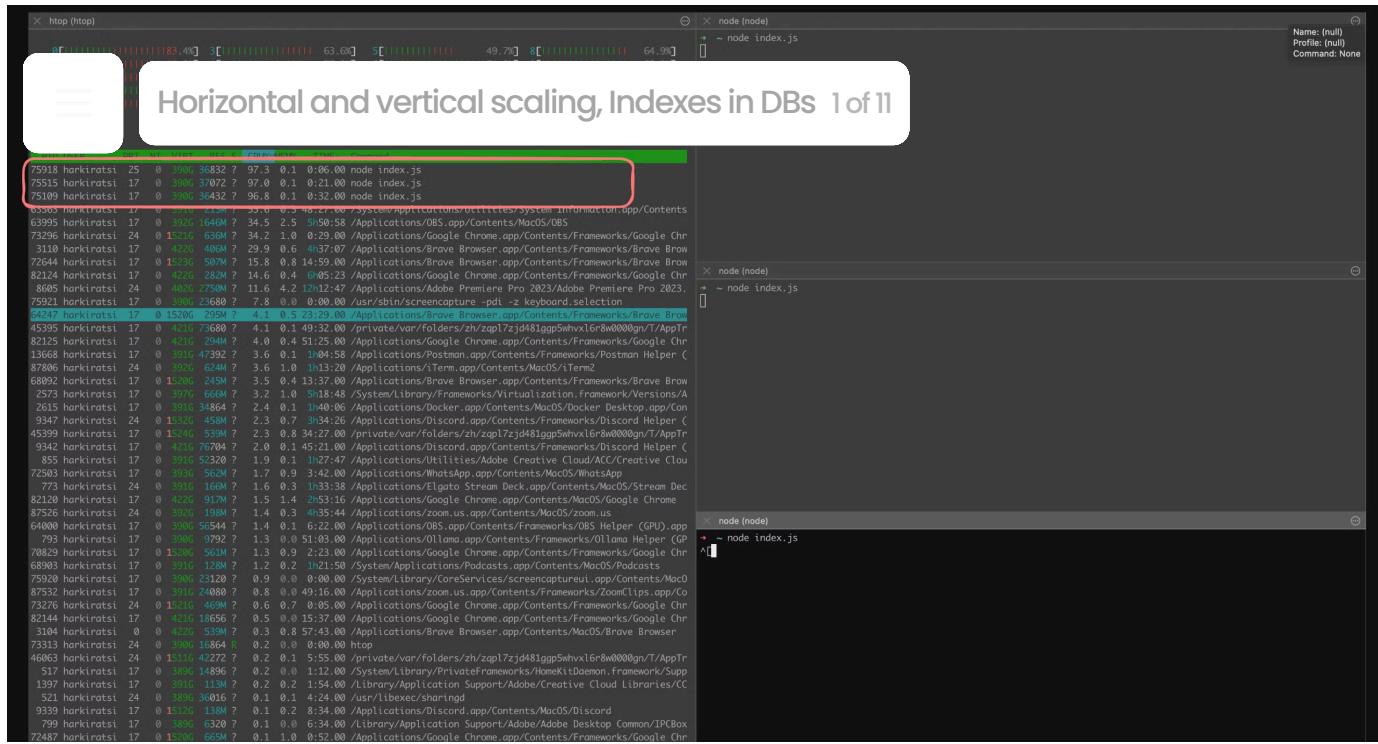
## Multi threaded languages



## Node.js

Let's run an infinite loop in a JS project and see how our CPU is used

```
let c = 0;  
while (1) {  
    c++;  
}
```



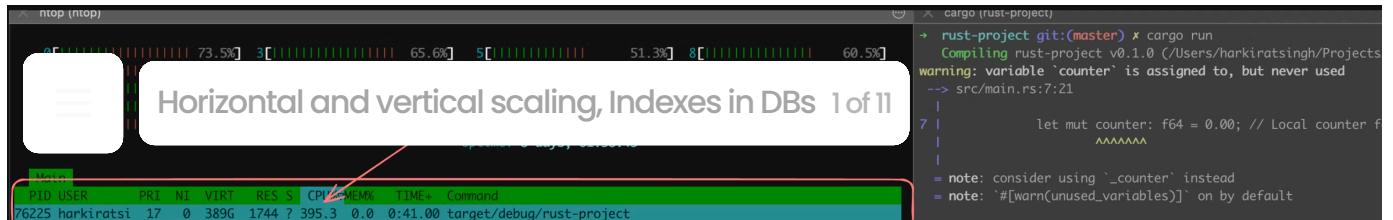
This confirms that only a single core of the machine is being used. We got 3 different processes using 100% CPU each.

## Rust

```
use std::thread;

fn main() {
    // Spawn three threads
    for _ in 0..3 {
        thread::spawn(|| {
            let mut counter: f64 = 0.00;
            loop {
                counter += 0.001;
            }
        });
    }

    loop {
        // Main thread does nothing but keep the program alive
    }
}
```



# Implementing horizontal scaling in Node.js project

You can start multiple node projects then? If there are 8 cores, then just start 8 projects?

```

node index.js

```

This, ofcourse has a lot of problems

1. Just ugly to do this, keep track of the processes that are up and down
2. Processes will have port conflicts, you'll have to run each process on a separate port

This is where the `cluster module` comes into the picture

```

import express from "express";
import cluster from "cluster";
import os from "os";

const totalCPUs = os.cpus().length;

```

```
Horizontal and vertical scaling, Indexes in DBs 1 of 11
console.log(`Primary ${process.pid} is running`);

// Fork workers.
for (let i = 0; i < totalCPUs; i++) {
  cluster.fork();
}

cluster.on("exit", (worker, code, signal) => {
  console.log(`worker ${worker.process.pid} died`);
  console.log("Let's fork another worker!");
  cluster.fork();
});
} else {
  const app = express();
  console.log(`Worker ${process.pid} started`);

  app.get("/", (req, res) => {
    res.send("Hello World!");
  });

  app.get("/api/:n", function (req, res) {
    let n = parseInt(req.params.n);
    let count = 0;

    if (n > 5000000000) n = 5000000000;

    for (let i = 0; i <= n; i++) {
      count += i;
    }

    res.send(`Final count is ${count} ${process.pid}`);
  });
}

app.listen(port, () => {
  console.log(`App listening on port ${port}`);
});
```

|

Browser

The screenshot shows a search result from Google. The title is "Horizontal and vertical scaling, Indexes in DBs 1 of 11". Below the title, there are several links and icons. One link is "https://getsubmis...". Other links include "hotel in paris - Go...", "ad", and "New". There are also icons for a gear, a play button, and a refresh.

Final count is 500500 78331

## Postman

The screenshot shows a Postman API request. The URL is "http://localhost:3000/api/1000". The method is "GET". The "Headers" tab is selected, showing the following headers:

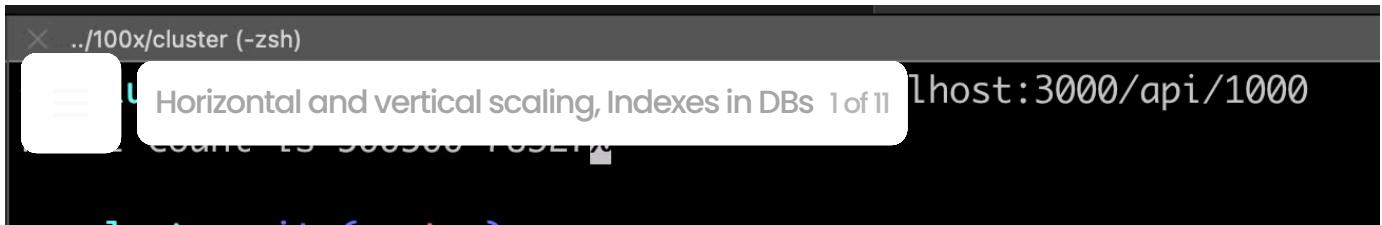
Key	Value
Accept	application/json, text/javascript, */*; q=0.01
Accept-Language	en-GB,en-US;q=0.9,en;q=0.8
Connection	keep-alive
Cookie	ci_session=0k0cgknftskse2bv5o7a5k0hb8c

The "Body" tab is selected, showing the response body:

```
1 Final count is 500500 78330
```

The status is 200 OK and the time is 4 ms.

## Curl



```
curl -sS localhost:3000/api/1000
```



Try to figure out why there is **stickiness** in the browser. Why the request from the same browser goes to the same pid

# Capacity estimation

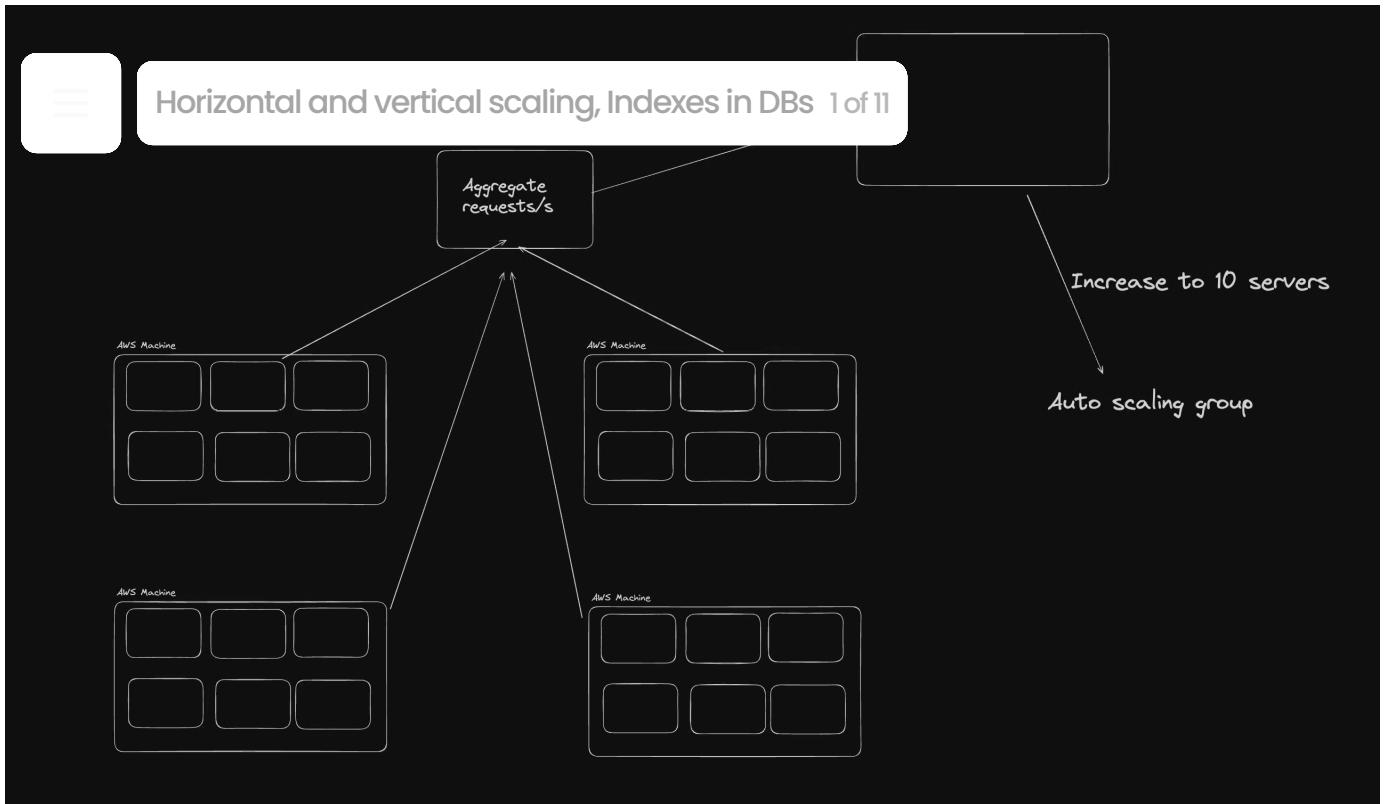
This is a common system design interview where they'll ask you

1. how would you scale your server
2. how do you handle spikes
3. How can you support a certain SLA given some traffic

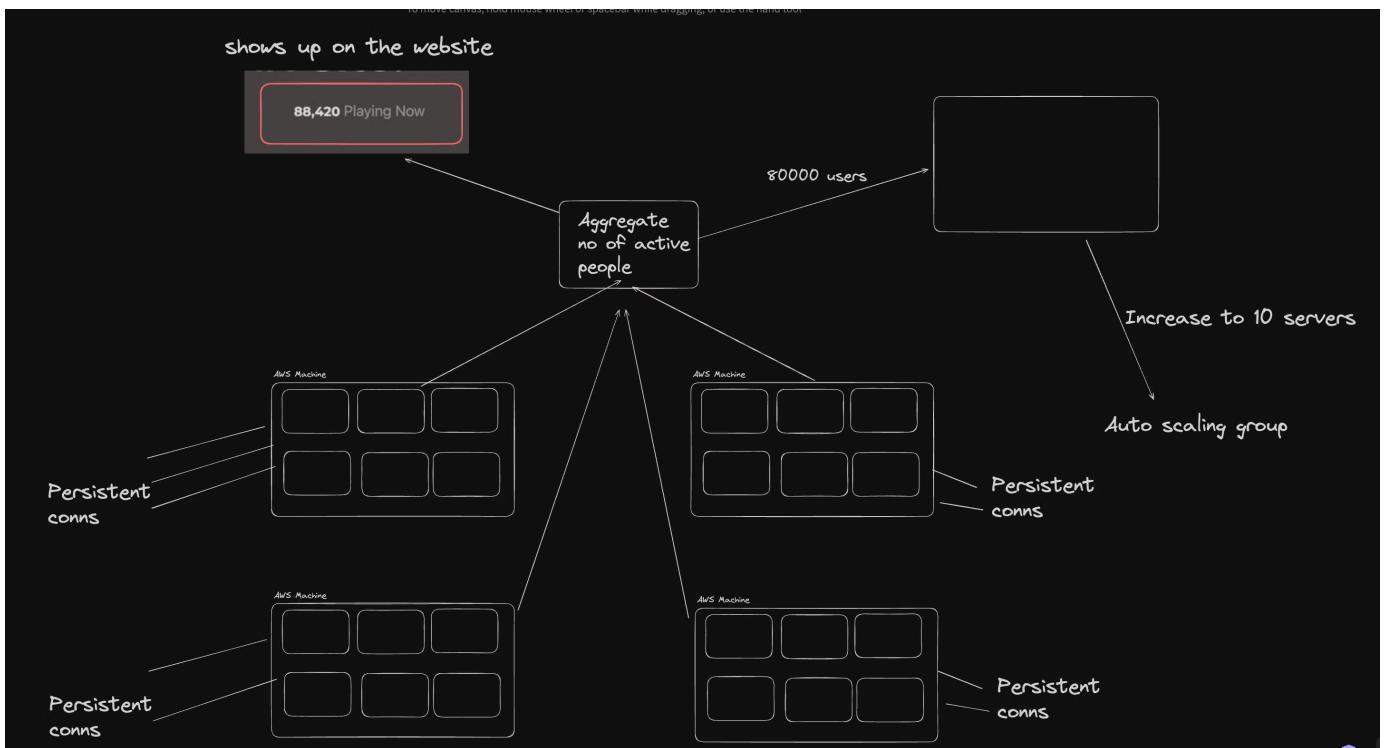
Answer usually requires a bunch of

1. paper math
2. Estimating requests/s
3. Assuming / monitoring how many requests a single machine can handle
4. Autoscaling machines based on the **load** that is estimated from time to time

## Example #1 – PayTM app



## Example #2 - Chess app



# Horizontal scaling

Horizontal and vertical scaling, Indexes in DBs 1 of 11

Horizontal scaling represents increasing the number of instances you have based on a metric to be able to support more load.

AWS has the concept of **Auto scaling groups**, which as the name suggests lets you autoscale the number of machines based on certain metrics.

## Buzz words

Images (AMI) - Snapshots of a machine from which you can create more machines

Load balancer - An entrypoint that your users send requests to that forwards it to one of many machines (or more specifically, to a target group). Its a **fully managed** service which means you don't have to worry about scaling it ever. AWS takes care of making it highly available.

Target groups - A group of EC2 instances that a load balancer can send requests to

Launch template - A template that can be used to start new machines



Please make sure you get rid of all your resources after this.

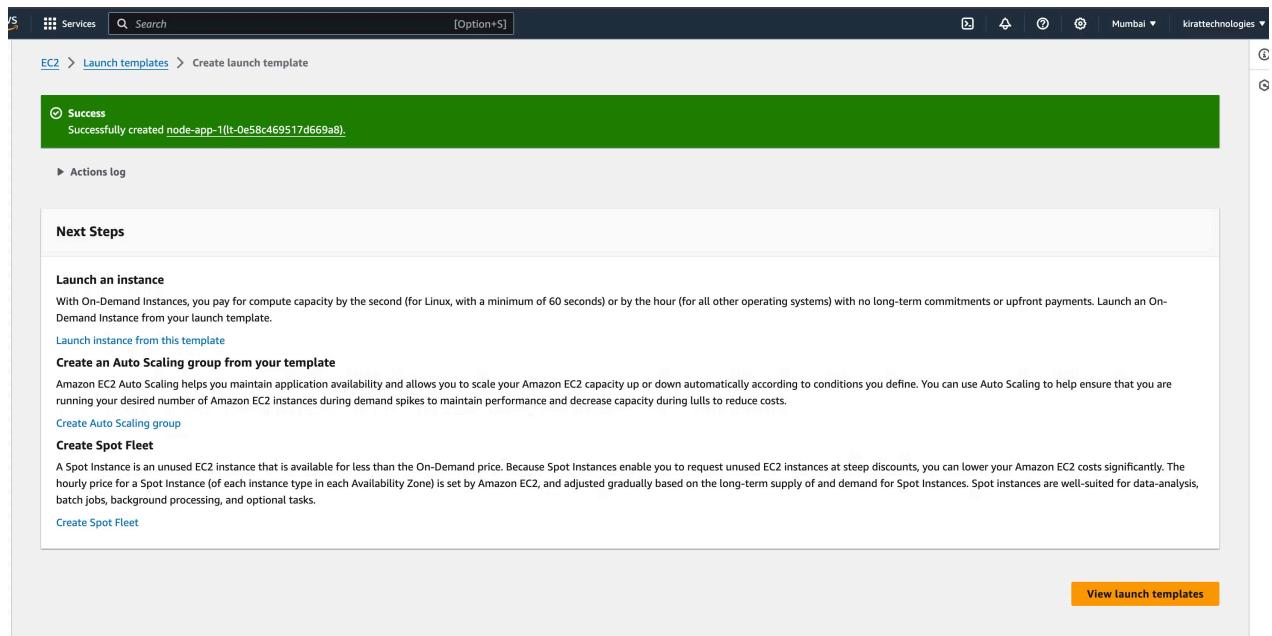
There are two ways you can use ASGs

- Create a EC2 instance.
  1. install Node.js on it  
<https://www.digitalocean.com/community/tutorials/how-to-install-node-js-on-ubuntu-20-04>
  2. Clone the repo - <https://github.com/100xdevs-cohort-2/week-22>
- Create an AMI with your machine

- Launch template

Horizontal and vertical scaling, Indexes in DBs 1 of 11  
<https://stackoverflow.com/questions/15904095/how-to-check-whether-my-user-data-passing-to-ec2-instance-is-working>

```
#!/bin/bash
export PATH=$PATH:/home/ubuntu/.nvm/versions/node/v22.0.0/bin/
echo "hi there before"
echo "hi there after"
npm install -g pm2
cd /home/ubuntu/week-22
pm2 start index.js
pm2 save
pm2 startup
```



- ASG

- Callout on availability zones – ASGs try to balance instances in each zone

## Horizontal and vertical scaling, Indexes in DBs 1 of 11

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

### VPC

Choose the VPC that defines the virtual network for your Auto Scaling group.

vpc-089cbcbf9c696895a  
172.31.0.0/16 Default



[Create a VPC](#)

### Availability Zones and subnets

Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets



ap-south-1a | subnet-0810906b79248efa6 X  
172.31.32.0/20 Default

ap-south-1b | subnet-04411ee701286912f X  
172.31.0.0/20 Default

[Create a subnet](#)

- Load balancer

- Add an HTTPS Listener from your domain, request a certificate from ACM

- Target group - Attach the target group to the ASG

## Autoscaling part

You can create an [dynamic scaling](#) policy

Horizontal and vertical scaling, Indexes in DBs 1 of 11

**Policy type**

Target tracking scaling

**Scaling policy name**

Target Tracking Policy

**Metric type** | [Info](#)  
Monitored metric that determines if resource utilization is too low or high. If using EC2 metrics, consider enabling detailed monitoring for better scaling performance.

Average CPU utilization

**Target value**

50

**Instance warmup** | [Info](#)  
300 seconds

Disable scale in to create only a scale-out policy

**Create**

Try playing with the Min and max on the ASG

node-app-1					
Details	Activity	Automatic scaling	Instance management	Monitoring	Instance refresh
<b>Group details</b>					<a href="#">Edit</a>
Auto Scaling group name node-app-1	Desired capacity 1	Desired capacity type Units (number of instances)	Status	Amazon Resource Name (ARN) <a href="#">arn:aws:autoscaling:ap-south-1:100749360009:autoScalingGroup:4ce2e6da-e011-44c8-8aa4-1f42afe1be0e:autoScalingGroupName/node-app-1</a>	
Date created Sun Apr 28 2024 15:01:23 GMT+0700 (Indochina Time)	Minimum capacity 1	Maximum capacity 4			
<b>Launch template</b>					<a href="#">Edit</a>

## Try killing servers

Try to stop a few servers in the ASG. Notice they spin back up to arrive at the desired amount.

Try running an infinite for loop on the instance to see if a scale up happens

Horizontal and vertical scaling, Indexes in DBs 1 of 11

```
while (1) {
    ++
}
```

```
ubuntu@ip-172-31-14-238: ~ (ssh)

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

Last login: Sun Apr 28 08:04:47 2024 from 101.96.67.218
To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-14-238:~$ ls
week-22
ubuntu@ip-172-31-14-238:~$ htop
ubuntu@ip-172-31-14-238:~$ vi index.js
ubuntu@ip-172-31-14-238:~$ node index.js
```

```
CPU[||||||||||||||||| 100.0%] Tasks: 44, 81 thr, 72 kthr; 1 running
Mem[||||||||||||||] 362M/1.92G Load average: 0.22 0.05 0.02
Swp[ 0K/0K] Uptime: 00:19:11

Main I/O
PID USER PRI NI VIRT RES SHR S CPU% MEM% TIME+ Command
1690 ubuntu 20 0 975M 46848 41216 R 99.3 2.3 0:11.10 node index.js
2013 ubuntu 20 0 8616 4224 3456 R 1.3 0.2 0:00.06 htop
    1 root 20 0 22244 13396 9556 S 0.0 0.7 0:05.68 /sbin/init
    138 root 20 0 66884 14636 13484 S 0.0 0.7 0:01.54 /usr/lib/systemd/s
    182 root RT 0 282M 27136 8704 S 0.0 1.3 0:00.03 /sbin/multipathd -
    185 root 20 0 26456 8180 5108 S 0.0 0.4 0:00.25 /usr/lib/systemd/s
    187 root 20 0 282M 27136 8704 S 0.0 1.3 0:00.00 /sbin/multipathd -
    188 root RT 0 282M 27136 8704 S 0.0 1.3 0:00.00 /sbin/multipathd -
    189 root RT 0 282M 27136 8704 S 0.0 1.3 0:00.00 /sbin/multipathd -
    190 root RT 0 282M 27136 8704 S 0.0 1.3 0:00.00 /sbin/multipathd -
    191 root RT 0 282M 27136 8704 S 0.0 1.3 0:00.07 /sbin/multipathd -
F1Help F2Setup F3Search F4Filter F5Tree F6SortBy F7Nice -F8Nice +F9Kill F10Quit
```

You'll notice the desired capacity goes up by one in some time

The screenshot shows a CloudTrail event log entry. The event details are as follows:

- Event Type:** 1 serv
- Instance ID:** i-011bf7ee8a3b91a34
- Description:** 4d15-0414-014017502cd in state ALARM triggered policy Target tracking Policy changing the desired capacity from 1 to 2. At 2024-04-28T08:27:29Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 1 to 2.
- Timestamp:** 2024 April 28, 03:27:31 PM +07:00

Try turning the infinite loop off and notice a scale down happens

## Scaling via a Node.js app

Create a new user with permissions to **AutoscalingFullAccess**

The left screenshot shows the AWS Lambda console. A red arrow points from the copied public IPv4 address (ec2-13-233-67-81.ap-south-1) to the right screenshot.

The right screenshot shows the AWS IAM User Management console. It displays the user's summary and the attached policies, which include the **AutoscalingFullAccess** policy.

```
import AWS from 'aws-sdk';
```

```
AWS.config.update({
  region: 'ap-south-1',
  accessKeyId: 'YOUR_ACCESS_KEY',
  secretAccessKey: 'YOUR_ACCESS_SECRET'
});
```

```
// Create an Auto Scaling client
const autoscaling = new AWS.AutoScaling();
```

```
// Function to update the desired capacity of an Auto Scaling group
const updateDesiredCapacity = (autoScalingGroupName: string, desiredCapacity: number) => {
  const params = {
    AutoScalingGroupName: autoScalingGroupName,
    DesiredCapacity: desiredCapacity
  };
  return autoscaling.setDesiredCapacity(params).promise();
}
```

```
autoscaling.setDesiredCapacity(params, (err, data) => {
  if (err) {
    Horizontal and vertical scaling, Indexes in DBs 1 of 11
  } else {
    console.log("Success", data);
  }
});

// Example usage
const groupName = 'node-app-1'; // Set your Auto Scaling group name
const newDesiredCapacity = 3; // Set the new desired capacity
```

# Indexing in Postgres

We've created postgres tables many times now. Let's see how/if indexing helps us speed up queries

- Create a postgres DB locally (don't use neon, we have a lot of data to store, will be very slow)

```
docker run -p 5432:5432 -e POSTGRES_PASSWORD=mysecretpassword -d postgres
```

- Connect to it and create some dummy data in it

```
docker exec -it container_id /bin/bash
psql -U postgres
```

- Create the schema for a simple medium like app

```
CREATE TABLE users (
  user_id SERIAL PRIMARY KEY,
  email VARCHAR(255) UNIQUE NOT NULL,
```

);

 RE Horizontal and vertical scaling, Indexes in DBs 1 of 11  
 p user\_id INTEGER NOT NULL,  
 title VARCHAR(255) NOT NULL,  
 description TEXT,  
 image VARCHAR(255),  
 FOREIGN KEY (user\_id) REFERENCES users(user\_id)  
 );

- Insert some dummy data in

```
DO $$  

DECLARE  

    returned_user_id INT;  

BEGIN  

    -- Insert 5 users  

    FOR i IN 1..5 LOOP  

        INSERT INTO users (email, password, name) VALUES  

        ('user'||i||'@example.com', 'pass'||i, 'User '||i)  

        RETURNING user_id INTO returned_user_id;  

        FOR j IN 1..500000 LOOP  

            INSERT INTO posts (user_id, title, description)  

            VALUES (returned_user_id, 'Title '||j, 'Description for post '||j);  

        END LOOP;  

    END LOOP;  

END $$;
```

- Try running a query to get all the posts of a user and log the time it took

```
EXPLAIN ANALYSE SELECT * FROM posts WHERE user_id=1 LIMIT 5;
```

Focus on the **execution time**

- Add an index to user\_id

```
CREATE INDEX idx_user_id ON posts (user_id);
```

Notice the **execution time** now.

What do you think happened that caused the query time to go down by so much?

## ☰ Horizontal and vertical scaling, Indexes in DBs 1 of 11

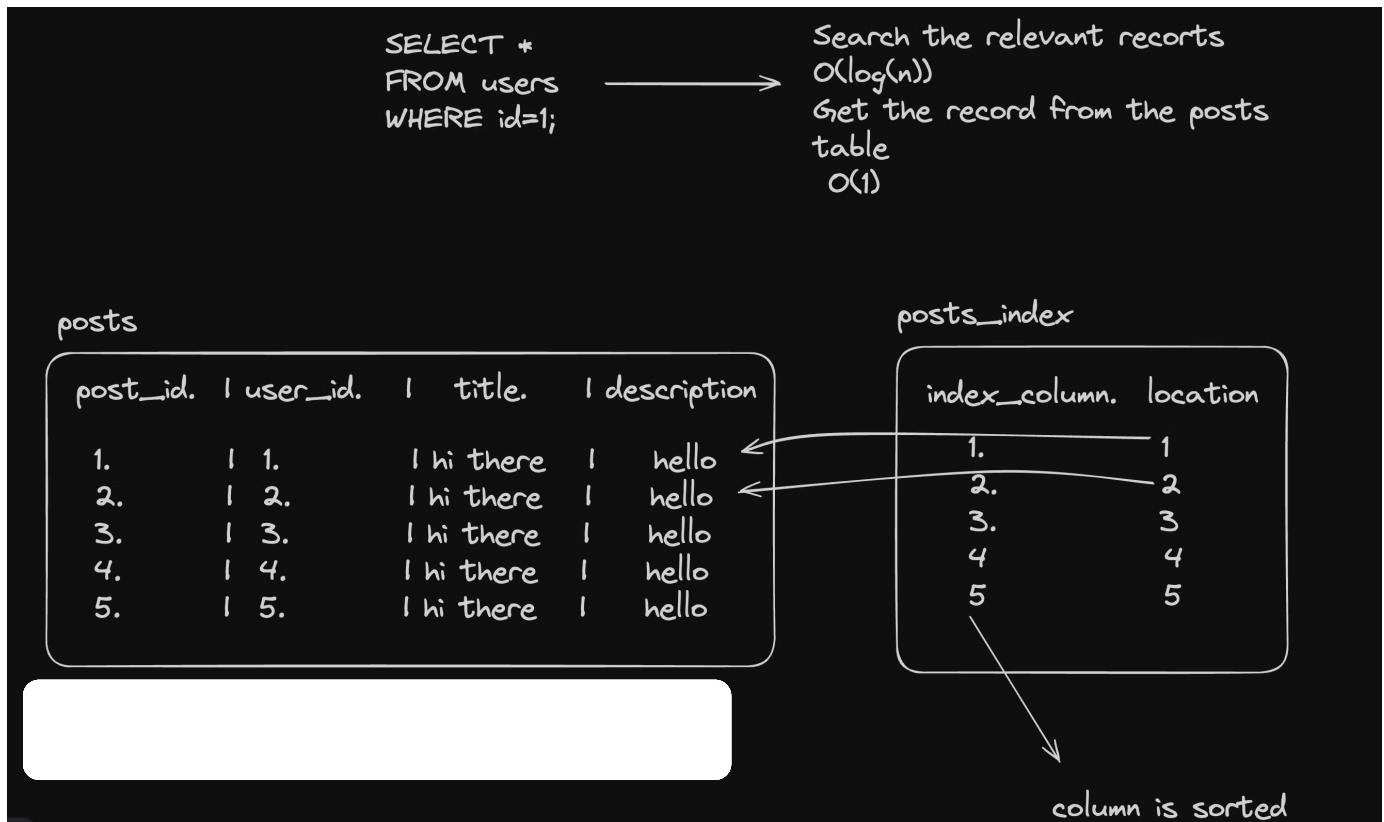
When you create an index on a field, a new data structure (usually B-tree) is created that stores the mapping from the **index column** to the **location** of the record in the original table.

Search on the index is usually  $\log(n)$

## Without indexes



## With indexes



The `data pointer` (in case of postares) is the `page` and `offset` at which this  
r | Horizontal and vertical scaling, Indexes in DBs 1 of 11

Think of the index as the `appendix` of a book and the `location` as the `page + offset` of where this data can be found

# Complex indexes

You can have index on more than one column for more complex queries

For example,

Give me all the posts of a user with given `id` with `title` "Class 1".

The index needs to have two keys now

```
CREATE INDEX idx_posts_user_id_title ON posts (description, title);
```



- Try searching before the index is added and after it is added

```
SELECT * FROM posts WHERE title='title' AND description='my title';
```



## Horizontal and vertical scaling, Indexes in DBs 1 of 11

# Indexes in Prisma

Ref - <https://www.prisma.io/docs/orm/prisma-schema/data-model/indexes>

You can add an index to a **model** in prisma by doing the following -

```
model User {  
    id      String @id @default(uuid())  
    username String @unique  
    email   String @unique  
    posts   Post[]  
    createdAt DateTime @default(now())  
    updatedAt DateTime @updatedAt  
}  
  
model Post {  
    id      String @id @default(uuid())  
    title   String  
    content String?  
    published Boolean @default(false)  
    createdAt DateTime @default(now())  
    updatedAt DateTime @updatedAt  
    userId   String  
    user     User  @relation(fields: [userId], references: [id])
```

{}

## Horizontal and vertical scaling, Indexes in DBs 1 of 11

Let's look at daily code and see where all can we introduce an index

<https://github.com/code100x/daily-code/blob/main/packages/db/prisma/schema.prisma#L129>

# Normalization

Normalization is the process of removing redundancy in your database.

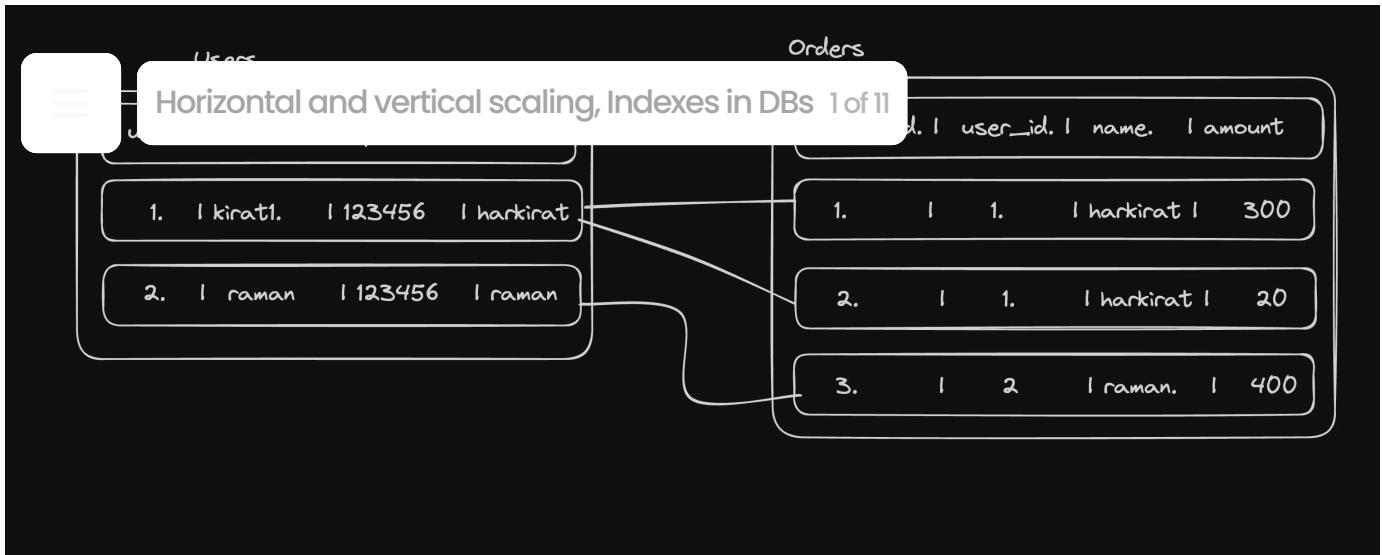
## Redundancy

Redundant data means data that already exists elsewhere and we're duplicating it in two places

For example, if you have two tables

1. users
2. user\_metadata

where you do the following -



If you notice, we've stored the name on the order in the Orders table, when it is already present in the Users table. This is what is **redundant** data.

Notice this schema is still **full proof**. We can get all the orders given a user id. We can tell the users details (username, name) given an order id.

## Non full proof data



This data doesn't have any relationship b/w Orders and users. This is just

## Types of relationships

## Use case - Library management system

☰ ↗ Horizontal and vertical scaling, Indexes in DBs 1 of 11

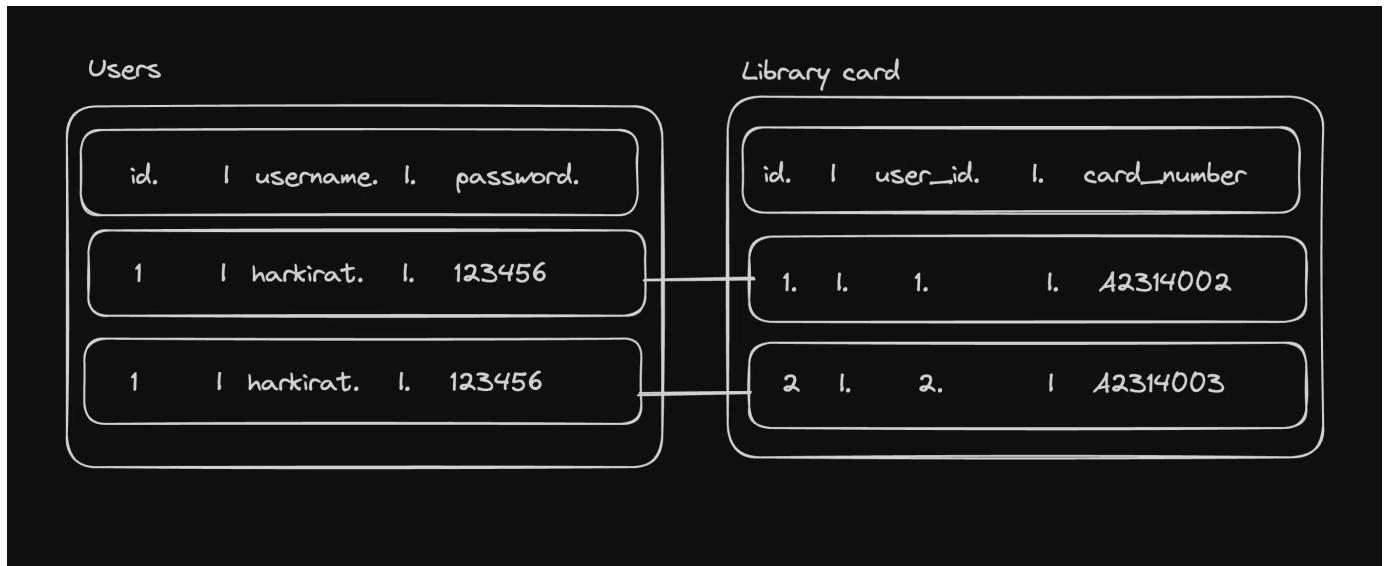
2. Library card table

3. Books table

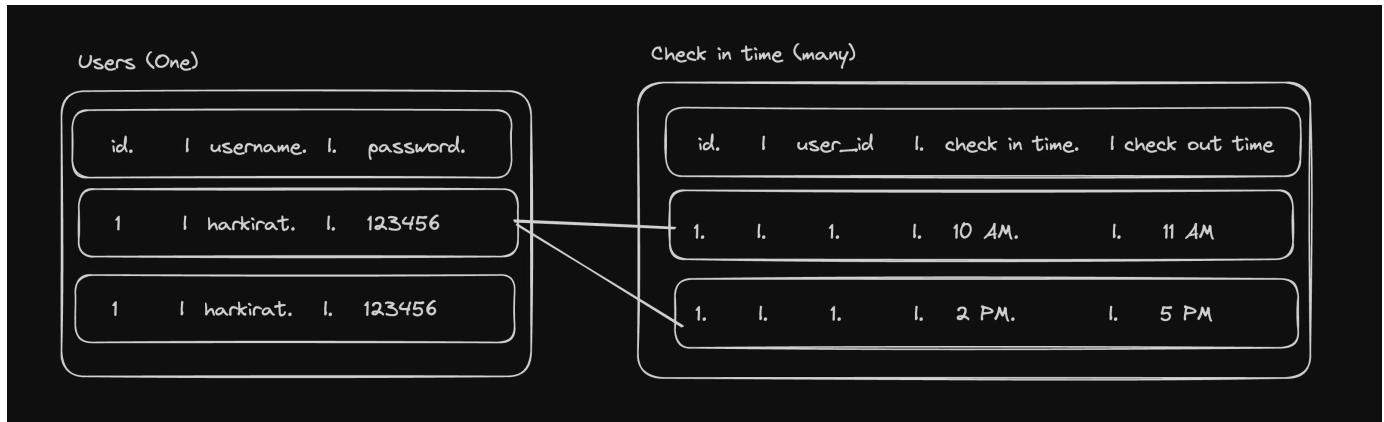
4. Genre table

## One to One

Each user has a single **Library card**

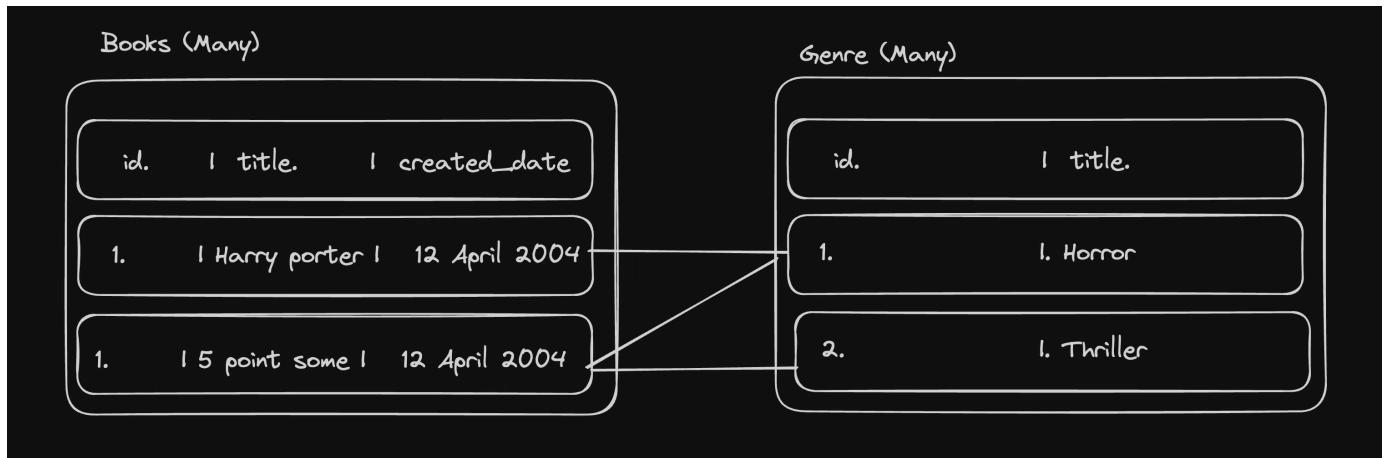


## One to many



Opposite of the thing above

## Horizontal and vertical scaling, Indexes in DBs 1 of 11



## Final graph

The diagram illustrates horizontal and vertical scaling of database indexes. It shows two separate tables. The first table has columns 'id.', 'title.', and 'created\_date'. A single row contains values 1., Harry porter, and 12 April 2004. The second table has columns 'id.' and 'title.'. A single row contains values 1. and Horror. This visualizes how data can be distributed across multiple tables (horizontal scaling) or how specific data like 'id.' can be replicated across rows (vertical scaling).

# Normalizing data

Normalization in databases is a systematic approach of decomposing tables to eliminate data redundancy and improve data integrity.

The process typically progresses through several normal forms, each building on the last.

When you look at a schema, you can identify if it lies in one of the following categories of normalization

1. 1NF
2. 2NF
3. 3NF
4. BCNF
5. 4NF
6. 5NF

You aim to reach 3NF/BCNF usually. The lower you go, the more normalised your table is. But over normalization can lead to excessive joins

## 1NF

- **A single cell must not hold more than one value (atomicity):** This rule ensures that each column of a database table holds only atomic (indivisible) values and multi-valued attributes are split into separate meant to store phone numbers, and

a person has multiple phone numbers, each number should be in a

### Horizontal and vertical scaling, Indexes in DBs 1 of 11

StudentID	Name	Activities
1	John Doe	Basketball, Soccer, Chess Club
2	Emily White	Drama Club; Yearbook

StudentID	Name	Activity
1	John Doe	Basketball
1	John Doe	Soccer
1	John Doe	Chess Club
2	Emily White	Drama Club
2	Emily White	Yearbook

- There must be a primary key for identification:** Each table should have a primary key, which is a column (or a set of columns) that uniquely identifies each row in a table
- No duplicated rows:** To ensure that the data in the table is organised properly and to uphold the integrity of the data, each row in the table should be unique. This rule works hand-in-hand with the presence of a primary key to prevent duplicate entries which can lead to data anomalies.
- Each column must have only one value for each row in the table:** This rule emphasizes that every column must hold only one value per row, and that value should be of the same kind for that column across all rows.

CustomerID	Name	ContactInfo
1	Alice Smith	alice@example.com, (555) 123-4567
2	Bob Johnson	bob@example.com; (555) 765-4321

CustomerID	Name	Email	Phone Number
1	Alice Smith	alice@example.com	(555) 123-4567
		bob@example.com	(555) 765-4321

## Horizontal and vertical scaling, Indexes in DBs 1 of 11

Ref - <https://www.studytonight.com/dbms/second-normal-form.php>

1NF gets rid of repeating rows. 2NF gets rid of redundancy

A table is said to be in 2NF if it meets the following criteria:

- is already in 1NF
- Has 0 partial dependency.

 Partial dependency - This occurs when a non-primary key attribute is dependent on part of a composite primary key, rather than on the whole primary key. In simpler terms, if your table has a primary key made up of multiple columns, a partial dependency exists if an attribute in the table is dependent only on a subset of those columns that form the primary key.

**Example:** Consider a table with the composite primary key (**StudentID**, **CourseID**) and other attributes like **InstructorName** and **CourseName**. If **CourseName** is dependent only on **CourseID** and not on the complete composite key (**StudentID**, **CourseID**), then **CourseName** has a partial dependency on the primary key. This violates 2NF.

## Before normalization

Enrollments table

StudentID	CourseID	CourseName	InstructorName	Grade
001	CS101	Computer Science	Dr. Smith	A
002	CS101	Computer Science	Dr. Smith	B
001	MA202	Mathematics	Dr. Johnson	A
003	CS101	Computer Science	Dr. Smith	C
002	MA202	Mathematics	Dr. Johnson	B

Can you spot the redundancy over here? The instructor name and course name should be the same for a given course!

Primary key of this table is (student\_id, course\_id)

CourseName and InstructorName have a **partial dependency** on CourseID

**Primary key**

StudentID	CourseID	CourseName	InstructorName	Grade
001	CS101	Computer Science	Dr. Smith	A
002	CS101	Computer Science	Dr. Smith	B
001	MA202	Mathematics	Dr. Johnson	A
003	CS101	Computer Science	Dr. Smith	C
002	MA202	Mathematics	Dr. Johnson	B

## After normalisation

CourseID	CourseName	InstructorName
CS101	Computer Science	Dr. Smith
MA202	Mathematics	Dr. Johnson

StudentID	CourseID	Grade
001	CS101	A
002	CS101	B
001	MA202	A
003	CS101	C
002	MA202	B

## 3NF

When a table is in 2NF, it eliminates repeating groups and redundancy, but it still has partial dependency.

So, for a table to be in 3NF, it must:

### Horizontal and vertical scaling, Indexes in DBs 1 of 11

- have no transitive partial dependency.



A **transitive dependency** in a relational database occurs when one non-key attribute indirectly depends on the primary key through another non-key attribute.

For example

Primary key				
EmployeeID	EmployeeName	DepartmentID	DepartmentName	DepartmentManager
1	Alice	D1	Sales	John
2	Bob	D2	Marketing	Sarah
3	Charlie	D1	Sales	John
4	Diana	D3	Finance	Mary

Department name has a **transitive dependency** on the primary key (employee id).

To normalise to 3NF, we need to do the following

EmployeeID	EmployeeName	DepartmentID
1	Alice	D1
2	Bob	D2
3	Charlie	D1
4	Diana	D3

DepartmentID	DepartmentName	DepartmentLocation
D1	Sales	New York
D2	Marketing	Los Angeles
D3	Finance	Chicago