



What we'll be doing today

What we've done already

1. Clusters
2. Nodes
3. Pods
4. Deployment
5. Replicaset

What we're doing today

1. Namespaces
2. Ingress
3. Ingress Controller
 1. nginx
 2. traefik
4. ConfigMaps
5. Secrets

What we're doing tomorrow

1. Cert management
2. Volumes and Persistent volumes





Online video next week

1. HPA - Horizontal Pod Autoscaling
2. Node autoscaling
3. Labs to add k8s a real codebase

Recapping what we've done

Ref - <https://projects.100xdevs.com/tracks/kubernetes-1/Kubernetes-Part-1-1>

Quick recap

▼ Cluster

A kubernetes cluster is a bunch of machines that work together to help you deploy your app

▼ Nodes

Each machine in your cluster is called a node

QUESTION

1 Master node (control plane) - Exposes an API that the **developer** can

☰ IS **Kubernetes Part 2** 1 of 23

2. Worker node - Actually run the **pods**

▼ Pods

The smallest execution unit inside a kubernetes cluseter. A pod can run **one or more** containers

▼ Replicaset

They let you create **multiple** pods (replicas).

It also takes care of bringing them back up if they ever go down/are killed

▼ Deployment

A deployment creates

▼ Services

Services let you expose your pods to other pods/over the internet



3. Loadbalancer – Creates a loadbalancer outside the kubernetes cluster

Nodeport

Loadbalancer

Recapping how to run this locally

Creating a cluster



nodes:

- role: control-plane
- extraPortMappings:
- containerPort: 30007
- hostPort: 30007
- role: worker
- extraPortMappings:
- containerPort: 30007
- hostPort: 30008
- role: worker

- Run the cluster locally

```
kind create cluster --config kind.yaml --name local2
```



- Run docker ps to confirm that the cluster is running

```
docker ps
```



Creating a pod

- Create a pod manifest (pod.yaml)

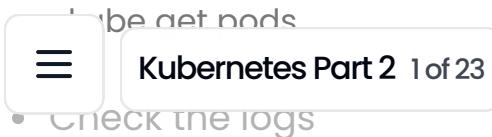
```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
```



- Apply the pod manifest

```
kubectl apply -f pod.yaml
```





kubectl logs -f nginx

- Delete the pod

kubectl **delete** pod nginx

Creating a replicaset

- Create the replicaset manifest

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
        ports:
          - containerPort: 80
```

- Apply the replicaset manifest

kubectl apply -f rs.yml

- Check the number of pods running now

- Try deleting a pod, and ensure it gets restarted



```
kubectl delete rs nginx-replicaset
```



Creating a Deployment

- Create a deployment manifest (deployment.yml)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
        ports:
          - containerPort: 80
```



- Apply the manifest

```
kubectl apply -f deployment.yml
```



- Check the `rs` that exist now

```
kubectl get rs
```



- Check the pods that exist now

```
curl http://127.0.0.1:80
```



- Try creating a new deployment with a wrong image name



```
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx2:latest
      ports:
        - containerPort: 80
```

- Ensure that the old pods are still running

```
kubectl get pods
```



Keep the deployment, it'll come in handy in the 4th slide

How to run this on a cloud

Go to a cloud provider like



Kubernetes Part 2 1 of 23

2. GCP

3. Digital ocean

4. Vultr

and create a k8s cluster

- Create a cluster

- Download the credentials file and replace `~/.kube/config` with it

- Create a deployment manifest

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
```

app: nginx

spec:



Kubernetes Part 2 1 of 23

image: nginx2:latest

ports:

- containerPort: 80

- Create a deployment

kubectl apply -f deployment.yml



Services

Services let you actually expose your app over the internet.

Nodeport

- Create a **Nodeport** service (service.yml)

apiVersion: v1



kind: Service

metadata:

name: nginx-service

spec:

selector:

app: nginx

ports:

- protocol: TCP

port: 80

targetPort: 80

nodePort: 30007 # This port needs to be any valid port within the NodePort range



- Apply it



- Visit any of the nodes on 30007

<http://localhost:30007/>



This will only work if you've started your `kind` cluster with the config from slide 2
On vultr, it will just work

LoadBalancer (will only work on a cloud provider)

The `LoadBalancer` service type is designed to work with cloud providers to create an external load balancer that routes traffic to the service.

- Replace the `type` to be `LoadBalancer`

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```



- Re-apply the config



```
kubectl apply -f service.yaml
```

Kubernetes Part 2 1 of 23



- See the loadbalancer created on the dashboard

- Visit the balancer to see the website

Downsides of services

Services are great, but they have some downsides -

Scaling to multiple apps

1. If you have three apps (frontend, backend, websocket server), you will have to create 3 separate services to route traffic to them. There is no way to do **centralized traffic management** (routing traffic from the same URL/Path Based Routing)

Load balancers you can create



Multiple certificates for every route

You can create certificates for your **load balancers** but you have to maintain them outside the cluster and create them manually

You also have to update them if they ever expire

No centralized logic to handle **rate limiting** to all services

Each load balancer can have its own set of rate limits, but you cant create a **single rate limitter** for all your services.

Here is a sample manifest that you can run to start two separate services to two separate LoadBalancer services

▼ Manifest

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:alpine
          ports:
            - containerPort: 80
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: apache-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: apache
  template:
    metadata:
      labels:
        app: apache
    spec:
```

Image: <https://raw.githubusercontent.com/k8s-samples/quick-start/1.23/deployments/deployment.yaml>

ports:



Kubernetes Part 2 1 of 23

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: LoadBalancer
```

```
apiVersion: v1
kind: Service
metadata:
  name: apache-service
spec:
  selector:
    app: apache
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: LoadBalancer
```

kubectl apply -f manifest.yml



You will notice two load balancers created for your two services



Ingress and Ingress Controller

Ref - <https://kubernetes.io/docs/concepts/services-networking/ingress/>

An API object that manages external access to the services in a cluster, typically HTTP.

Ingress may provide load balancing, SSL termination and name-based virtual hosting.



An Ingress does not expose arbitrary ports or protocols. Exposing services other than HTTP and HTTPS to the internet typically uses a service of

service.Type=LoadBalancer.



Ingress controller

If you remember from last week, our `control plane` had a `controller manager` running.

Ref - <https://projects.100xdevs.com/tracks/kubernetes-1/Kubernetes-Part-1-3>

The `kube-controller-manager` runs a bunch of `controllers` like

1. Replicaset controller
2. Deployment controller

etc

If you want to expose your services outside of the Kubernetes cluster, you need to install an `Ingress Controller`. These come by default in k8s

Famous k8s ingress controllers



Kubernetes Part 2 1 of 23

ller for Kubernetes works with the NGINX webserver (as a proxy).

- HAProxy Ingress is an ingress controller for HAProxy.
- The Traefik Kubernetes Ingress provider is an ingress controller for the Traefik proxy.

Full list – <https://kubernetes.io/docs/concepts/services-networking/ingress-controllers/>

Namespaces

In Kubernetes, a **namespace** is a way to divide cluster resources between multiple users/teams. Namespaces are intended for use in environments with many users spread across multiple teams, or projects, or environments like development, staging, and production.

When you do

```
kubectl get pods
```



it gets you the **pods** in the **default** namespace

Creating a new namespace

- Create a new namespace

```
curl -H "Content-Type: application/json" -d @create-namespace.json http://127.0.0.1:8080/api/v1/namespaces
```



- Get all the namespaces



Kubernetes Part 2 1 of 23

b



- Get all pods in the namespace

```
kubectl get pods -n my-namespace
```



- Create the manifest for a deployment in the namespace

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  namespace: backend-team
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
```



- Apply the manifest

```
kubectl apply -f deployment-ns.yml
```



- Get the deployments in the namespace

```
kubectl get deployment -n backend-team
```



- Get the pods in the namespace

```
kubectl get pods -n backend-team
```



- Set the default context to be the namespace



Kubernetes Part 2 1 of 23

b

--current --namespace=backend-team



- Try seeing the pods now

kubectl get pods



- Revert back the kubectl config

kubectl config set-context --current --namespace=default



Install the nginx ingress controller

Ref - <https://docs.nginx.com/nginx-ingress-controller/installation/installing-nic/installation-with-manifests/>

Using helm

- Install helm

Ref - <https://helm.sh/>

Installation - <https://helm.sh/docs/intro/install/>

- Add the **ingress-nginx** chart

helm repo add ingress-nginx <https://kubernetes.github.io/ingress-nginx>



helm repo update

helm install nginx-ingress ingress-nginx/ingress-nginx --namespace ingress-



- Check if you have pods running in the



Default loadbalancer service

You will notice that if you use `helm` to install the `nginx-ingress-controller`, it creates a `Loadbalancer` service for you

```
kubectl get services --all-namespaces
```



This routes all the traffic to an `nginx pod`

```
kubectl get pods -n ingress-nginx
```



This means the first part of our `ingress deployment` is already created



Adding the routing to the ingress controller

Next up, we want to do the following -

- Get rid of all existing deployments in the default namespace

```
kubectl get deployments
```



```
kubectl delete deployment_name
```

- Get rid of all the services in the default namespace (dont delete the default kubernetes service, delete the old `nginx` and `apache` loadbalancer services)

```
kubectl get services
```



```
kubect
```

- Create a `deployment` and `service` definition for the `nginx` image/app (this is different from the nginx controller)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  namespace: default
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
```



app: nginx



Kubernetes Part 2 1 of 23

```
- name: nginx  
  image: nginx:alpine  
  ports:  
    - containerPort: 80
```

```
apiVersion: v1  
kind: Service  
metadata:  
  name: nginx-service  
  namespace: default  
spec:  
  selector:  
    app: nginx  
  ports:  
    - protocol: TCP  
      port: 80  
      targetPort: 80  
  type: ClusterIP
```

- Create a deployment and service for the apache app

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: apache-deployment  
  namespace: default  
spec:  
  replicas: 2  
  selector:  
    matchLabels:  
      app: apache  
template:  
  metadata:  
    labels:  
      app: apache  
spec:  
  containers:
```

ports.

```
- containerPort: 80
```



Kubernetes Part 2 1 of 23

```
kind: Service
metadata:
  name: apache-service
  namespace: default
spec:
  selector:
    app: apache
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: ClusterIP
```

- Create the ingress resource

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: web-apps-ingress
  namespace: default
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  rules:
    - host: your-domain.com
      http:
        paths:
          - path: /nginx
            pathType: Prefix
            backend:
              service:
                name: nginx-service
                port:
                  number: 80
          - path: /apache
            pathType: Prefix
            backend:
```

name: apache-service

Combined manifest

- ▼ Create a combined manifest with all the api objects

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: nginx-deployment
```

```
  namespace: default
```

```
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

```
      app: nginx
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: nginx
```

```
    spec:
```

```
      containers:
```

```
        - name: nginx
```

```
          image: nginx:alpine
```

```
          ports:
```

```
            - containerPort: 80
```

```
---
```

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: nginx-service
```

```
  namespace: default
```

```
spec:
```

```
  selector:
```

```
    app: nginx
```

```
  ports:
```

```
    - protocol: TCP
```

```
      port: 80
```

```
---
```

apiVersion: apps/v1
kind: Deployment
name: apache-deployment
namespace: default
spec:
 replicas: 2
 selector:
 matchLabels:
 app: apache
 template:
 metadata:
 labels:
 app: apache
 spec:
 containers:
 - name: my-apache-site
 image: httpd:2.4
 ports:
 - containerPort: 80

```
apiVersion: v1
kind: Service
metadata:
  name: apache-service
  namespace: default
spec:
  selector:
    app: apache
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: ClusterIP
```

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: web-apps-ingress
  namespace: default
  annotations:
    kubernetes.io/ingress.class: "nginx"
spec:
  rules:
    - host: "web-apps-ingress.default.svc.cluster.local"
      http:
        paths:
          - path: /
            backend:
              service:
                name: apache-service
                port:
                  number: 80
```

ingressClassName: nginx



Kubernetes Part 2 1 of 23 m

```
http:  
  paths:  
    - path: /nginx  
      pathType: Prefix  
      backend:  
        service:  
          name: nginx-service  
          port:  
            number: 80  
    - path: /apache  
      pathType: Prefix  
      backend:  
        service:  
          name: apache-service  
          port:  
            number: 80
```

- Apply the manifest

kubectl apply -f complete.yml



- Update your local hosts entry (`/etc/hosts`) such that `your-domain.com` points to the IP of your load balancer

65.20.84.86 your-domain.com



- Try going to `your-domain.com/apache` and `your-domain.com/nginx`

Trying traefik's ingress controller

Traefik is another popular ingress controller. Let's try to our apps using it next

- Install traefik ingress controller using helm

```
helm repo add traefik https://helm.traefik.io/traefik
```



```
helm repo update
```

```
helm install traefik traefi/traefik --namespace traefik --create-namespace
```

- Make sure an ingressClass is created for traefik

```
kubectl get IngressClass
```



- Notice it created a **LoadBalancer** svc for you

```
kubectl get svc -n traefik
```



- Create a **Ingress** that uses the traefik ingressClass and traefik annotations (traefik.yml)

```
apiVersion: networking.k8s.io/v1
```



```
kind: Ingress
```

```
metadata:
```

```
  name: traefik-web-apps-ingress
```

```
  namespace: default
```

```
spec:
```

```
  ingressClassName: traefik
```

```
  rules:
```

```
    host: traefik.domain.com
```

paths.

- path: /nginx



Kubernetes Part 2 1 of 23

service:

 name: nginx-service

 port:

 number: 80

- path: /apache

 pathType: Prefix

 backend:

 service:

 name: apache-service

 port:

 number: 80

- Add an entry to your `/etc/hosts` (IP should be your loadbalancer IP)

65.20.90.183 traefik-domain.com



- Visit the website

traefik-domain.com/nginx



traefik-domain.com/apache



Can you guess what is going wrong? Why are you not seeing anything on this final page?

Assignment

Try to figure out how can you rewrite the path to `/` if you're using traefik as

t



Secrets and configmaps

Kubernetes suggests some standard configuration practises.

These include things like

1. You should always create a deployment rather than creating naked pods
2. Write your configuration files using YAML rather than JSON
3. Configuration files should be stored in version control before being pushed to the cluster

Kubernetes v1 API also gives you a way to store **configuration** of your application outside the image/pod

This is done using

1. ConfigMaps
2. Secrets

Rule of thumb

Don't store secrets or configuration in your docker image

Pass them in as environment variables whenever you're starting the



ConfigMaps

Ref - <https://kubernetes.io/docs/concepts/configuration/configmap/>

A ConfigMap is an API object used to store non-confidential data in key-value pairs. Pods can consume ConfigMaps as environment variables, command-line arguments, or as configuration files in a volume.

A ConfigMap allows you to decouple environment-specific configuration from your container images, so that your applications are easily portable.

Creating a ConfigMap

- Create the manifest

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: ecom-backend-config
data:
  database_url: "mysql://ecom-db:3306/shop"
  cache_size: "1000"
```



ecom-gateway.ecommerce.svc.cluster.local

ecom-gateway.example.com

max_cart_items: "50"



e

Kubernetes Part 2 1 of 23

- Apply the manifest

```
kubectl apply -f cm.yaml
```



- Get the configmap

```
kubectl describe configmap ecom-backend-config
```



Creating an express app that exposes env variables

▼ Express app code

```
import express from 'express';
import fs from 'fs';
import path from 'path';

const app = express();
const port = 3000;
app.get('/', (req, res) => {
  const envVars = {
    DATABASE_URL: process.env.DATABASE_URL,
    CACHE_SIZE: process.env.CACHE_SIZE,
    PAYMENT_GATEWAY_URL: process.env.PAYMENT_GATEWAY_URL,
    MAX_CART_ITEMS: process.env.MAX_CART_ITEMS,
    SESSION_TIMEOUT: process.env.SESSION_TIMEOUT,
  };

  res.send(`

<h1>Environment Variables</h1>
<pre>${JSON.stringify(envVars, null, 2)}</pre>
`);

  app.listen(port, () => {
    console.log(`App listening at http://localhost:${port}`);
  });
});
```



FROM node:20



Kubernetes Part 2 1 of 23

WORKDIR /usr/src/app

```
COPY package*.json ./  
RUN npm install
```

```
COPY ..
```

```
RUN npx tsc -b
```

```
EXPOSE 3000  
CMD [ "node", "index.js" ]
```

▼ Deploy to dockerhub –

<https://hub.docker.com/repository/docker/100xdevs/env-backend/general>

Trying the express app using docker locally

```
docker run -p 3003:3000 -e DATABASE_URL=asd 100xdevs/env-backend
```



Try running using k8s locally

- Create the manifest (express-app.yml)

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: ecom-backend-deployment
```



```
  selector:
```

matchLabels:



Kubernetes Part 2 1 of 23

metadata:

labels:

app: ecom-backend

spec:

containers:

- name: ecom-backend
image: 100xdevs/env-backend

ports:

- containerPort: 3000

env:

- name: DATABASE_URL

valueFrom:

configMapKeyRef:

name: ecom-backend-config
key: database_url

- name: CACHE_SIZE

valueFrom:

configMapKeyRef:

name: ecom-backend-config
key: cache_size

- name: PAYMENT_GATEWAY_URL

valueFrom:

configMapKeyRef:

name: ecom-backend-config
key: payment_gateway_url

- name: MAX_CART_ITEMS

valueFrom:

configMapKeyRef:

name: ecom-backend-config
key: max_cart_items

- name: SESSION_TIMEOUT

valueFrom:

configMapKeyRef:

name: ecom-backend-config
key: session_timeout

- Apply the manifest

https://kubernetes.io/docs/tutorials/kubernetes-basics/deployment/deployment-manage-applications/



- Create the service (express-service.yml)



```
kind: Service
metadata:
  name: ecom-backend-service
spec:
  type: NodePort
  selector:
    app: ecom-backend
  ports:
    - port: 3000
      targetPort: 3000
      nodePort: 30007
```

- Apply the service

```
kubectl apply -f express-service.yml
```



- Try visiting the website

Secrets

Secrets are also part of the kubernetes v1 api. They let you store **passwords** / **sensitive data** which can then be mounted on to pods as environment variables. Using a Secret means that you don't need to include confidential

C



Using a secret

- Create the manifest with a secret and pod (secret value is base64 encoded) (<https://www.base64encode.org/>)

```
apiVersion: v1
kind: Secret
metadata:
  name: dotfile-secret
data:
  .env: REFUQUJBU0VfVVJMPSJwb3N0Z3JlczovL3VzZXJuYWlIOnNIY3JldEBsb2Nhbi
---
apiVersion: v1
kind: Pod
metadata:
  name: secret-dotfiles-pod
spec:
  containers:
    - name: dotfile-test-container
      image: nginx
      volumeMounts:
        - name: env-file
          readOnly: true
          mountPath: "/etc/secret-volume"
  volumes:
    - name: env-file
      secret:
        secretName: dotfile-secret
```



ing the .env

```
! kubectl exec -it secret-dotfiles-pod /bin/bash
```

Kubernetes Part 2 1 of 23



Base64 encoding

Whenever you're storing values in a secret, you need to base64 encode them. They can still be **decoded**, and hence this is not for security purposes. This is more to provide a standard way to store secrets, incase they are binary in nature.

For example, TLS (https) certificates that we'll be storing as secrets eventually can have non ascii characters. Converting them to base64 converts them to ascii characters.

Secrets as env variables

You can also pass in secrets as environment variables to your process (similar to how we did it for configmaps in the last slide)

- Create the secret

```
apiVersion: v1
kind: Secret
metadata:
  name: my-secret
data:
  username: YWRtaW4= # base64 encoded 'admin'
  password: cGFzc3dvcmQ= # base64 encoded 'password'
```



- Create the pod

```
apiVersion: v1
kind: Pod
metadata:
  name: secret-env-pod
spec:
```



```
image: busybox
  command: ["sh", "-c", "echo Username: $USERNAME; echo Password: $PASSWORD"]
  - name: USERNAME
    valueFrom:
      secretKeyRef:
        name: my-secret
        key: username
  - name: PASSWORD
    valueFrom:
      secretKeyRef:
        name: my-secret
        key: password
```

ConfigMaps vs Secrets

- Creating a ConfigMap

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: example-config
data:
  key1: value1
  key2: value2
```

- Creating a Secret

```
apiVersion: v1
kind: Secret
metadata:
  name: example-secret
```



Key differences

- **Purpose and Usage:**

- **Secrets:** Designed specifically to store sensitive data such as passwords, OAuth tokens, and SSH keys.
- **ConfigMaps:** Used to store non-sensitive configuration data, such as configuration files, environment variables, or command-line arguments.

- **Base64 Encoding:**

- **Secrets:** The data stored in Secrets is base64 encoded. This is not encryption but simply encoding, making it slightly obfuscated. This encoding allows the data to be safely transmitted as part of JSON or YAML files.
- **ConfigMaps:** Data in ConfigMaps is stored as plain text without any encoding.

- **Volatility and Updates:**

- **Secrets:** Often, the data in Secrets needs to be rotated or updated more frequently due to its sensitive nature.
- **ConfigMaps:** Configuration data typically changes less frequently compared to sensitive data.

- **Kubernetes Features:**

- **Secrets:** Kubernetes provides integration with external secret management systems and supports encryption at rest for Secrets when configured properly. Ref <https://secrets-store-csi-driver.sigs.k8s.io/concepts.html#provider-for-the-secrets-store-csi-driver>
- **ConfigMaps:** While ConfigMaps are used to inject configuration data into containers, they do not provide the same level of support for external



Adding https using cert-manager

Ref - <https://cert-manager.io/>

Try installing a certificate for a domain name of your own before tomorrow's class

Maybe get a domain name from namecheap for cheap -
<https://www.namecheap.com/>



Volumes in docker

Pretext

The following docker image runs a Node.js app that writes periodically to the filesystem -

<https://hub.docker.com/r/100xdevs/write-random>

▼ Nodejs Code

```
const fs = require('fs');
const path = require('path');

// Function to generate random data
function generateRandomData(length) {
    let characters = 'ABCDEFGHIJKLMNOPQRSTUVWXYZabcdefghijklmnopqrstuvwxyz';
    let result = "";
    for (let i = 0; i < length; i++) {
        result += characters.charAt(Math.floor(Math.random() * characters.length));
    }
    return result;
}

// Write random data to a file
function writeRandomDataToFile(filePath, dataLength) {
    const data = generateRandomData(dataLength);
    fs.writeFile(filePath, data, (err) => {
        if (err) {
            console.error(`Error writing file: ${err}`);
        } else {
            console.log(`File written successfully`);
        }
    });
}
```

```
    console.log('Data written to file', filePath);
```



Kubernetes Part 2 1 of 23

```
}
```

```
// Define the file path and data length
const filePath = path.join(__dirname, '/generated/randomData.txt');
const dataLength = 100; // Change this to desired length of random data

// Write random data to file every 10 seconds
setInterval(() => {
  writeRandomDataToFile(filePath, dataLength);
}, 10000); // 10000 ms = 10 seconds

// Keep the script running
console.log('Node.js app is running and writing random data to randomDat
```

Run it in docker

Try running the image above in your local machine

```
docker run 100xdevs/write-random
```



Try going to the container and seeing the contents of the container

```
docker exec -it container_id /bin/bash
cat randomData.txt
```



Where is this file being stored?

The data is stored in the **docker runtime filesystem**. When the container dies, the data dies with it. This is called **ephemeral storage**

Volumes in docker





If you want to persist data across container stops and starts, you can use Volumes in Docker

Bind mounts

Replace the mount on the left with a folder on your own machine

```
docker run -v /Users/harkiratsingh/Projects/100x/mount:/usr/src/app/generat
```



Volume Mounts

- Create a volume

```
docker volume create hello
```



- Mount data to volume

```
docker run -v hello:/usr/src/app/generated 100xdevs/write-random
```



If you stop the container in either case, the `randomFile.txt` file persists

Volumes in kubernetes

Ref - <https://kubernetes.io/docs/concepts/storage/volumes/>



In Kubernetes, a Volume is a directory, possibly with some data in it, which is part of its filesystem. Kubernetes supports a variety of volume types, such as EmptyDir, PersistentVolumeClaim, Secret, ConfigMap, and others.

Why do you need volumes?

- If two containers in the same `pod` want to share data/fs.
- If you want to create a database that persists data **even when** a container restarts (creating a DB)
- Your `pod` just needs extra space during execution (for caching lets say) but doesn't care if it persists or not.

Types of volumes

Temporary volume that can be shared amongst various containers of a

 Kubernetes Part 2 1 of 23 e volume dies with it.

For example -

1. ConfigMap
2. Secret
3. emptyDir

Persistent Volume

A Persistent Volume (PV) is a piece of storage in the cluster that has been provisioned by an administrator or dynamically provisioned using Storage Classes. It is a resource in the cluster just like a node is a cluster resource. PVs are volume plugins like Volumes but have a lifecycle independent of any individual Pod that uses the PV. This API object captures the details of the implementation of the storage, be that NFS, iSCSI, or a cloud-provider-specific storage system.

Persistent volume claim

A Persistent Volume Claim (PVC) is a request for storage by a user. It is similar to a Pod. Pods consume node resources and PVCs consume PV resources. Pods can request specific levels of resources (CPU and Memory). Claims can request specific size and access modes (e.g., can be mounted once read/write or many times read-only).

Ephemeral volumes

A temporary volume that can be shared amongst various containers of a pod to share data.
Ephemeral volumes die with the pod.



Setup

- Create a manifest that starts two pods which share the same volume

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: shared-volume-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: shared-volume-app
  template:
    metadata:
      labels:
        app: shared-volume-app
    spec:
      containers:
        - name: writer
          image: busybox
          command: ["/bin/sh", "-c", "echo 'Hello from Writer Pod' > /data/hello.txt; sleep 3600"]
        volumeMounts:
          - name: shared-data
            mountPath: /data
        - name: reader
          image: busybox
          command: ["/bin/sh", "-c", "cat /data/hello.txt; sleep 3600"]
        volumeMounts:
          - name: shared-data
            mountPath: /data
      volumes:
        - name: shared-data
```

- Apply the manifest



- Check the reader container and see if you can see the volume data in there

```
kubectl exec -it shared-volume-deployment-74d67d6567-tcdsl --container
```



Persistent volumes

Just like our kubernetes cluster has **nodes** where we provision our **pods**.

We can create **persistent volumes** where our **pods** can **claim** (ask for) storage

Persistent volumes can be provisioned statically or dynamically.



Static persistent volumes

Creating a NFS

NFS is one famous implementation you can use to deploy your own persistent volume

I'm running one on my aws server -

```
version: '3.7'
```



```
services:
```

```
nfs-server:
```

```
  image: itsthenetwork/nfs-server-alpine:latest
```

```
  container_name: nfs-server
```

```
  privileged: true
```

```
  environment:
```

```
    SHARED_DIRECTORY: /exports
```

```
  volumes:
```

```
    - ./data:/exports:rw
```

- "2049:2049"



Kubernetes Part 2 1 of 23



Make sure the 2049 port on your machine is open

Creating a pv and pvc

Create a persistent volume claim and persistent volume

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: nfs-pv
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteMany
  storageClassName: nfs
  nfs:
    path: /exports
    server: 52.66.197.168
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: nfs-pvc
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 10Gi
  storageClassName: nfs
```

Create a pod

```
apiVersion: v1
kind: Pod
metadata:
```

name: mongo-pod



Kubernetes Part 2 1 of 23

```
- name: mongo
  image: mongo:4.4
  command: ["mongod", "--bind_ip_all"]
  ports:
    - containerPort: 27017
  volumeMounts:
    - mountPath: "/data/db"
      name: nfs-volume
  volumes:
    - name: nfs-volume
  persistentVolumeClaim:
    claimName: nfs-pvc
```

Try it out

- Put some data in mongodb

```
kubectl exec -it mongo-pod -- mongo
use mydb
db.mycollection.insert({ name: "Test", value: "This is a test" })
exit
```



- Delete and restart the pod

```
kubectl delete pod mongo-pod
kubectl apply -f mongo.yml
```



- Check if the data persists

```
kubectl exec -it mongo-pod -- mongo
use mydb
db.mycollection.find()
```



Automatic pv creation

Ref - <https://docs.vultr.com/how-to-provision-persistent-volume-claims-on-vultr-kubernetes-engine>

- Create a persistent volume claim with storageClassName set to `vultr-block-storage-hdd`

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: csi-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 40Gi
  storageClassName: vultr-block-storage-hdd
```

- Apply the pod manifest

```
apiVersion: v1
kind: Pod
metadata:
  name: mongo-pod
spec:
  containers:
    - name: mongo
      image: mongo:4.4
      command: ["mongod", "--bind_ip_all"]
      ports:
        - containerPort: 27017
  volumes:
```

VOLUMEMENTS

- name: mongo-storage



Kubernetes Part 2 1 of 23

- name: mongo-storage
persistentVolumeClaim:
claimName: csi-pvc

- Explore the resources created

kubectl get pv



kubectl get pvc

kubectl get pods

- Put some data in mongodb

kubectl exec -it mongo-pod -- mongo



use mydb

db.mycollection.insert({ name: "Test", value: "This is a test" })

exit

- Delete and restart the pod

kubectl **delete** pod mongo-pod



kubectl apply -f mongo.yml

- Check if the data persists

kubectl exec -it mongo-pod -- mongo



use mydb

db.mycollection.find()

